


# Understanding Exception- Related Bugs in Large-Scale Cloud Systems

ASE'19

Haicheng Chen, Wensheng Dou  
Yanyan Jiang, Feng Qin



Haicheng Chen  
Ohio State University

[ASE'19] **Understanding Exception-Related Bugs in Large-Scale Cloud Systems** [p]  
*Haicheng Chen, Wensheng Dou, Yanyan Jiang, Feng Qin*

[FSE'16] **Crash Consistency Validation Made Easy**  
*Yanyan Jiang, Haicheng Chen, Feng Qin, Chang Xu, Xiaoxing Ma, Jian Lu*



Feng Qin  
Ohio State University

**2019**

ASE

highly dependable and secure computer systems

**Understanding Exception-Related Bugs in Large-Scale Cloud Systems**

Haicheng Chen, Wensheng Dou, Yanyan Jiang, and Feng Qin

*In Proceedings of the 34th IEEE/ACM International Conference on Automated Software Engineering*

USENIX ATC

**Lessons and Actions: What We Learned from 10K SSD-Related Storage System Failures**

Erci Xu, Mai Zheng, Feng Qin, Yikang Xu, and Jiesheng Wu

*In Proceedings of the 2019 USENIX Annual Technical Conference*, Jul. 2019

**2018**

PDSW-DISCS

**Understanding SSD Reliability in Large-Scale Cloud Systems**

Erci Xu, Mai Zheng, Feng Qin, Yikang Xu, and Jiesheng Wu

*In Proceedings of the 3rd ACM/IEEE Joint International Workshop on Parallel Data Storage and Analysis*, Nov. 2018

FSE

**Distinguished  
Paper**

**An Empirical Study on Crash Recovery Bugs in Large-Scale Distributed Systems**


Yu Gao, Wensheng Dou, Feng Qin, Chushu Gao, Dong Wang, Jun Wei, Ruirui Huang, Li Z

*In Proceedings of the 26th ACM Joint European Software Engineering Conference and Symposium on Foundations of Software Engineering*

# Motivation

- Exception mechanism is **widely used** in cloud systems
  - **7%** of the source code in twelve popular open source distributed systems
- The **sophisticated logic** of cloud systems hinder use of exception mechanism
- Mistakes in the exception mechanism use may lead to **severe consequences**(eBugs)
  - system downtime, data loss.....

# Exception Mechanism

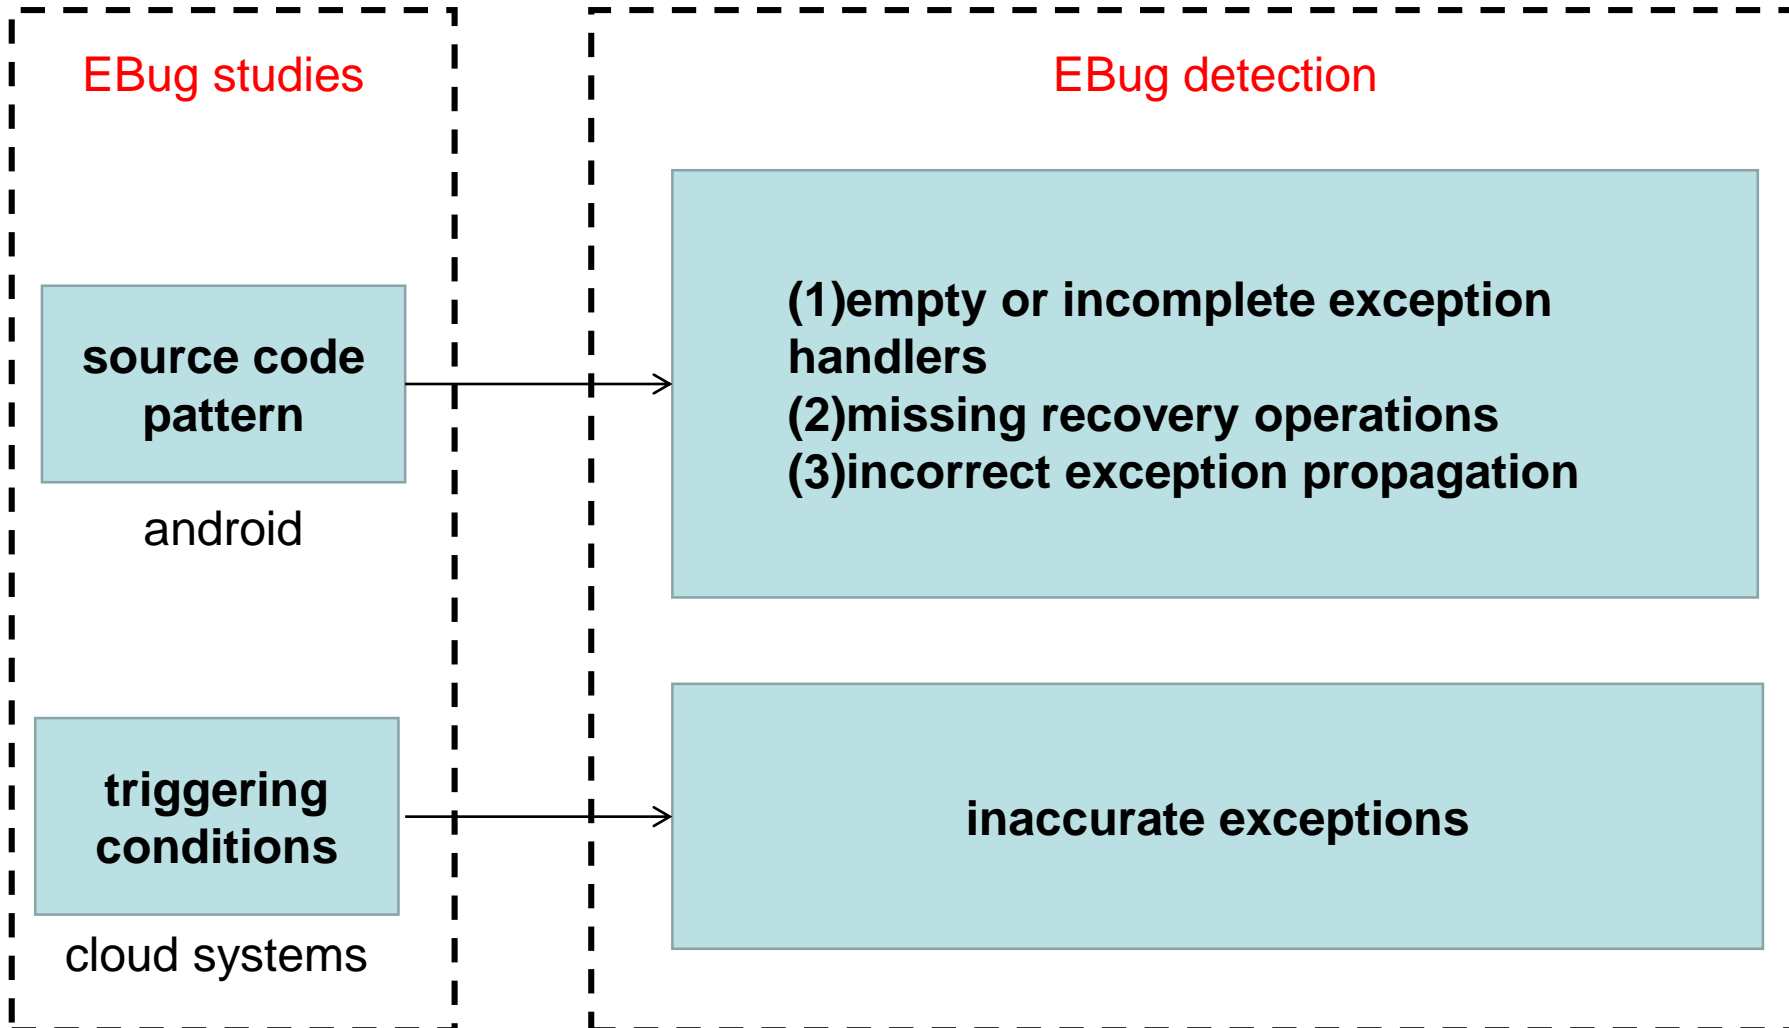


```
void B(...) throws OtherException {  
    try { A(...);  
    } catch (SomeException e) {  
        someHandling(...);  
        throw new OtherException(e);  
    }  
}
```

# Contributions

- Present the first comprehensive **study** on eBugs from the perspective of **triggering conditions**
- Build a static analysis **tool**, called DIET, and evaluate it using the latest versions of the studied systems
- Provide a large **benchmark of eBugs** in cloud systems

# Related Work



# Study on eBugs

- Data source: **JIRA**

System	Cassandra	HBase	HDFS	MR	YARN	ZK	Total
Retrieved	1,336	1,576	763	460	457	210	4,802
Studied	40	92	31	16	23	8	210

- eBug Analysis
  - RQ1: **Triggering conditions** of eBugs
  - RQ2: **Relation** between the **triggering conditions** and the **root causes** of eBugs
  - RQ3: **The impacts** of eBugs on cloud systems

# Triggering Condition Types

Triggering Condition (# eBugs)		Scenario (# eBugs)	CA	HB	HF	MR	YN	ZK
Non-semantic condition (114)	Network error (46)	Premature disconnection (17)	0	8	5	1	1	2
		Local timeout (12)	2	5	2	0	3	0
		Connection refused (11)	1	8	0	1	1	0
		Other network errors (6)	0	6	0	0	0	0
	File system error (40)	File corrupted (23)	10	11	1	1	0	0
		File not found (13)	3	5	3	2	0	0
		Other file system errors (4)	2	0	2	0	0	0
	Out of resource (16)	Out of memory (5)	1	2	1	1	0	0
		Out of disk space (5)	1	0	1	1	1	1
		Port conflicted (3)	1	1	1	0	0	0
		Out of other resources (3)	0	0	3	0	0	0
	Untimely interrupt (12)	Thread interrupted when invoking a blocking method (12)	2	1	0	3	6	0
Semantic condition (96)	-	-	17	45	12	6	11	5
<b>Total</b>	<b>210</b>		<b>40</b>	<b>92</b>	<b>31</b>	<b>16</b>	<b>23</b>	<b>8</b>



# Timing Requirements on Triggering Conditions

- **Weak:** occur at any consistent global state
  - ZOOKEEPER-2757 can be triggered whenever a user issues a delete command with **an invalid pathname**
- **Moderate:** occur on a node when it is in certain states
- **Strong:** occur on a node when both the current node and other nodes are in certain states

# Timing Requirements on Triggering Conditions

Condition Type	Timing Requirement		
	Weak	Moderate	Strong
Network error	9	36	1
File system error	32	5	3
Out of resource	5	11	0
Untimely interrupt	0	10	2
Semantic condition	39	34	23
<b>Total</b>	<b>85</b>	<b>96</b>	<b>29</b>

Most (86%) of the eBugs do not have strong timing requirements on their triggering conditions.

# Root Causes

- Inaccurate Exception
  - **creating** an inaccurate exception
- Missing Reaction
  - neither **catching** nor **specifying** an exception in the method signature
- Overly-General Reaction
  - different exceptions are incorrectly **handled** in the same way
- Incorrect Reaction Logic
  - a **handler** is incorrect for all the exceptions it handles

# Root Causes

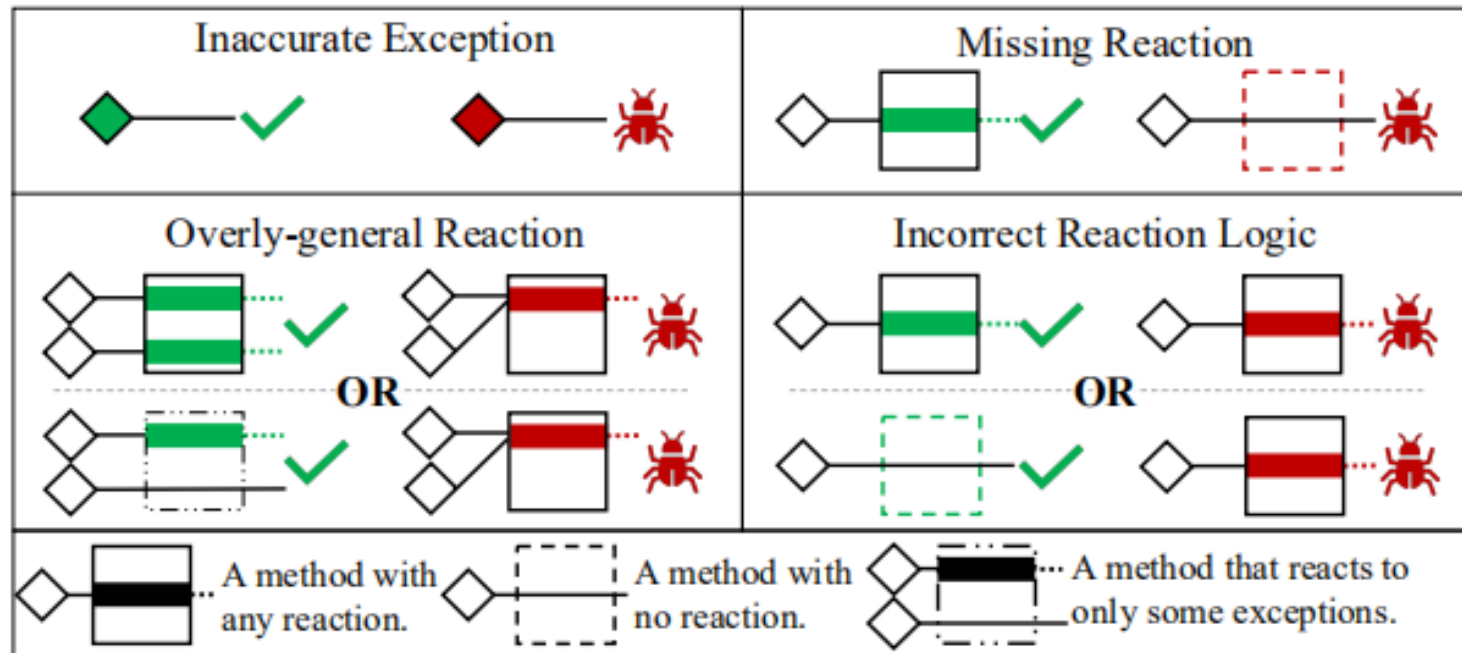


Fig. 3. Four different types of eBug root causes. For each type, we show the correct version on the left (green) and the buggy version on the right (red).

# Root Causes

Root Cause	eBug #	CA	HB	HF	MR	YN	ZK
Inaccurate exception	21	3	8	4	3	3	0
Missing reaction	36	12	11	3	3	4	3
Overly-general reaction	87	13	42	14	6	11	1
Incorrect reaction logic	66	12	31	10	4	5	4
<b>Total</b>	<b>210</b>	<b>40</b>	<b>92</b>	<b>31</b>	<b>16</b>	<b>23</b>	<b>8</b>

# Relation between triggering conditions and root causes

- Inaccurate Exception

Type	Wrong Class	Wrong Message	Lacking Cause	Total
eBug #	13	5	3	21

```
void updateMetaLocation() throws IOException {
    if (waitForRootServerConnection() == null)
-   throw new NullPointerException(...);
+   throw new IOException(...);
}
void process() {
    try { updateMetaLocation();
    } catch (IOException e) { cleanup(); }
}
```

In half of the eBugs that create a totally misleading exception, the exception class is **inconsistent** with its triggering condition.

# Relation between triggering conditions and root causes

- Overly-General Reaction

Relation	Same Type	Different Types	Unknown	Total
eBug #	48	30	9	87

- In **many (34%)** overly-general reaction eBugs, the incorrectly reacted exception and the correctly reacted ones are caused by **different types** of triggering conditions.

# Bug Impacts

Symptom	eBug #
Node downtime	48
Incorrect error message	44
Data loss or potential data loss	31
Hang or performance downgrading	26
Resource leak/exhaustion	10
Operation failure <sup>†</sup>	51
<b>Total</b>	<b>210</b>

Priority	Blocker	Critical	Major	Minor	Trivial	Total
eBug #	21	42	110	33	4	<b>210</b>



# Detecting eBugs in Cloud Systems

- Inaccurate exceptions
  - checking if the exceptions are consistent with their triggering conditions.
- Overly-general reactions
  - checking if the handled exceptions are triggered by different condition types.

# DIET: Detecting Inaccurate Exceptions

the root exception and triggering conditions

$$P_{c,t} = \frac{\text{number of } (c_i, t_j) \text{ where } c_i = c, t_i = t}{\text{number of } (c_i, t_i) \text{ where } c_i = c} \quad (1)$$

error message triggering conditions

$$P_{w,t} = \frac{\text{number of } (w_{i,j}, t_i) \text{ where } w_{i,j} = w, t_i = t}{\text{number of } (w_{i,j}, t_i) \text{ where } w_{i,j} = w} \quad (2)$$

$$P_{m,t} = \frac{\sum_{w_i \in m} P_{w_i,t}}{n}$$

Detect

$$P_{\text{same-type}} = \sum_{t \in \text{Five types}} \min(P_{c,t}, P_{m,t}) \quad P_{\text{same-type}} \leq 0.2$$

# DIET: Detecting Inaccurate Exceptions

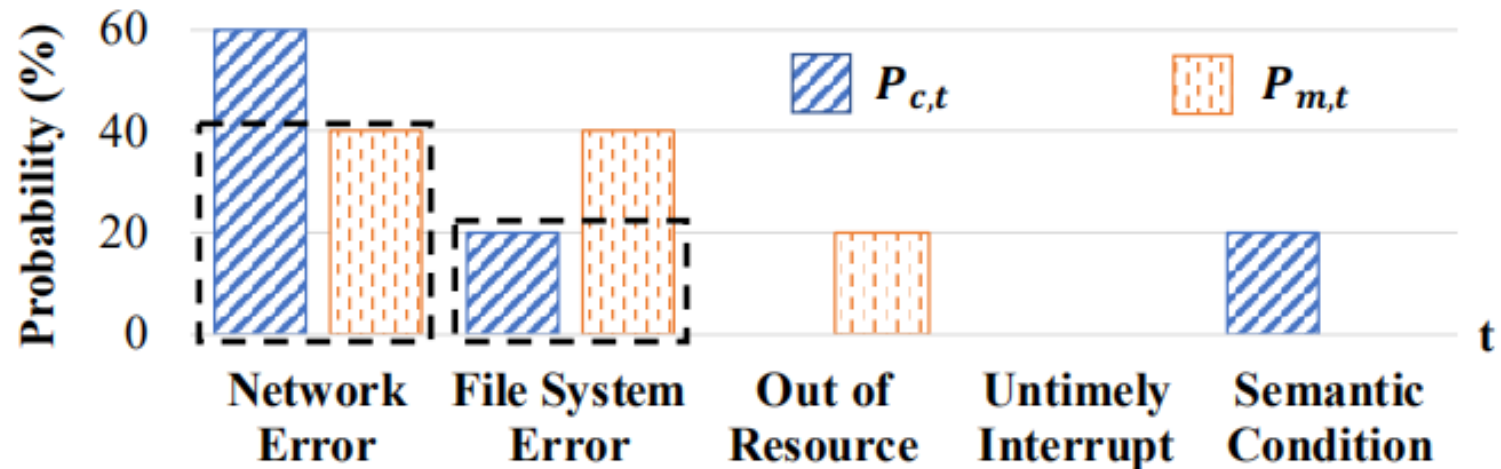


Fig. 8. The  $P_{c,t}$  and  $P_{m,t}$  of an exception. For example, when  $t$  is network error,  $P_{c,t}$  is 60% and  $P_{m,t}$  is 40%. The overlapping areas are highlighted with the dashed boxes. The total overlap is 60%.

# Result

## APPLYING DIET ON REAL-WORLD CLOUD SYSTEMS

<b>System (Version)</b>	Cassandra (3.11)	Hadoop <sup>†</sup> (3.1.2)	HBase (2.1.4)	ZooKeeper (2.4.14)	<b>Total</b>
<b>Throw</b>	2,823	9,853	5,020	429	<b>18,125</b>
<b>Root ex.</b>	1,282	3,090	1,374	159	<b>5,905</b>
<b>Calculated</b>	550	1,579	716	84	<b>2,929</b>
<b>Reported</b>	100	136	73	5	<b>314</b>
<b>Candidate</b>	9	20	2	0	<b>31</b>