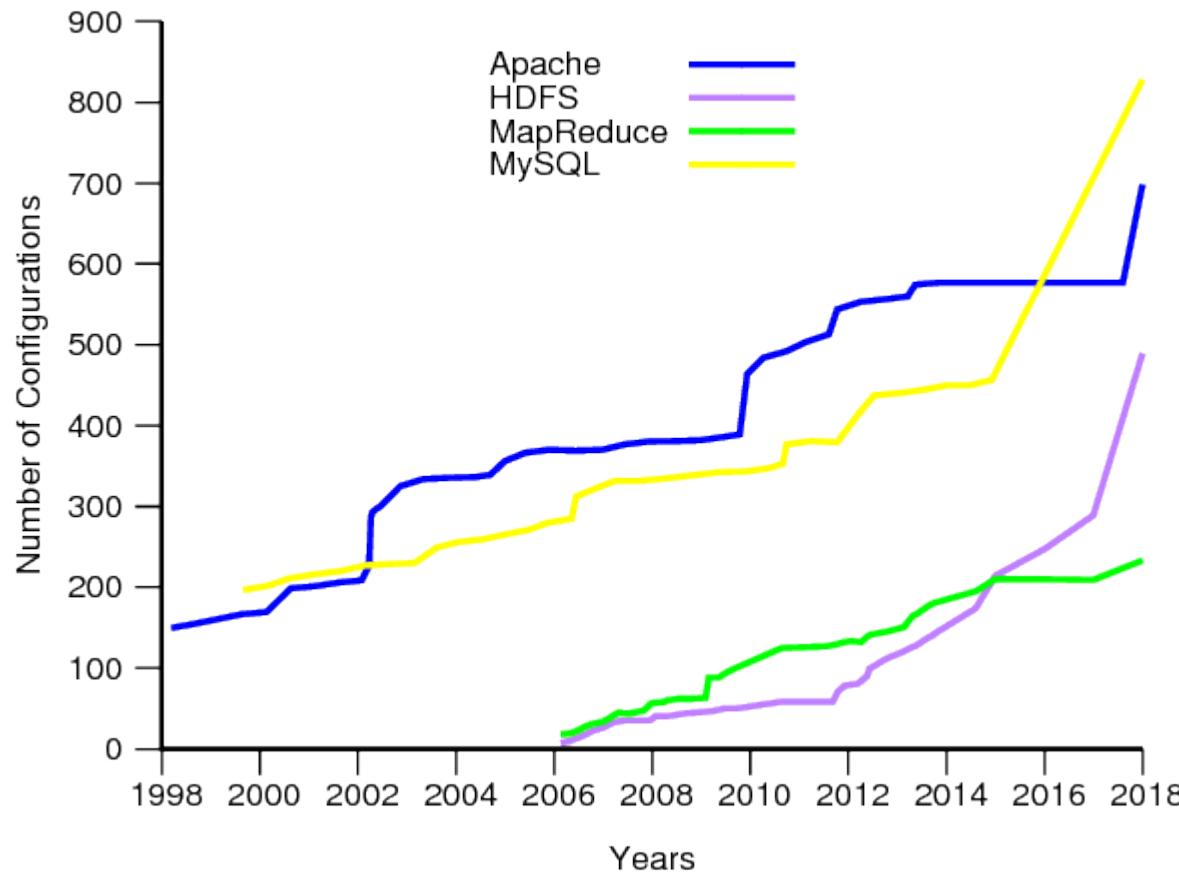


# Statically Inferring Performance Properties of Software Configurations

**Chi Li, Shu Wang, Henry Hoffmann, Shan Lu**



# Configurations Explosion



Tianyin Xu, Long Jin, Xuepeng Fan, Yuanyuan Zhou, Shankar Pasupathy, and Rukma Talwadker . Hey, You Have Given Me Too Many Knobs! Understanding and Dealing with Over-Designed Configuration in System Software. In FSE, 2016

# Which configuration affects performance?



stack**overflow** #36170959, Cassandra Performance Tuning

*"Please let me know **what more settings I can tweak** to get maximum performance out of my cluster."*



stack**overflow** #47665640, Memory configurations

*"I am finding that I am running out of memory when running my queries. I was able to figure out how to restrict cassandra to run in less than 4gb. **Is there such a setting for hadoop?**"*



stack**overflow** #45565896, MapReduce Error: Java heap space

*"Besides those parameters in the configuration, I do not change anything else, so I use the default values. **How can I solve the Error: Java Heap Space**"*

# How to performance-tune configurations?



stack**overflow** #37897438, Hbase Performance Tuning

*“I have the following parameters in Hbase: ... Can anyone suggest any configuration changes to generate more IO per second?”*



stack**overflow** #7243670, Hbase performance

*“My major configurations are: ... Am I doing something wrong with the configuration? This is my last shot at Hbase. Please help”*



Jira #HBase-13919, Rationalize Client Timeout

*“There are currently many settings that influence how/when an HBase client times out. This is hard to configure, hard to understand, and badly documented.”*

# Performance Misconfigurations

- Common
  - 65% of configuration issue reports
  - 35% of configuration posts on Stack Overflow
- Severe
  - 20% of MySQL misconfig. -> severe slowdown
  - 1/3 of Hadoop misconfig. -> memory issue (**OOM**)

Shu Wang, Chi Li, Henry Hoffmann, Shan Lu, William Sentosa, and Achmad Imam Kistijantoro. Understanding and auto-adjusting performance-sensitive configurations. In ACM SIGPLAN Notices, volume 53, pages 154–168. ACM, 2018.

Zuoning Yin, Xiao Ma, Jing Zheng, Yuanyuan Zhou, Lakshmi N Bairavasundaram, and Shankar Pasupathy. An empirical study on configuration errors in commercial and open source systems. In SOSP, 2011

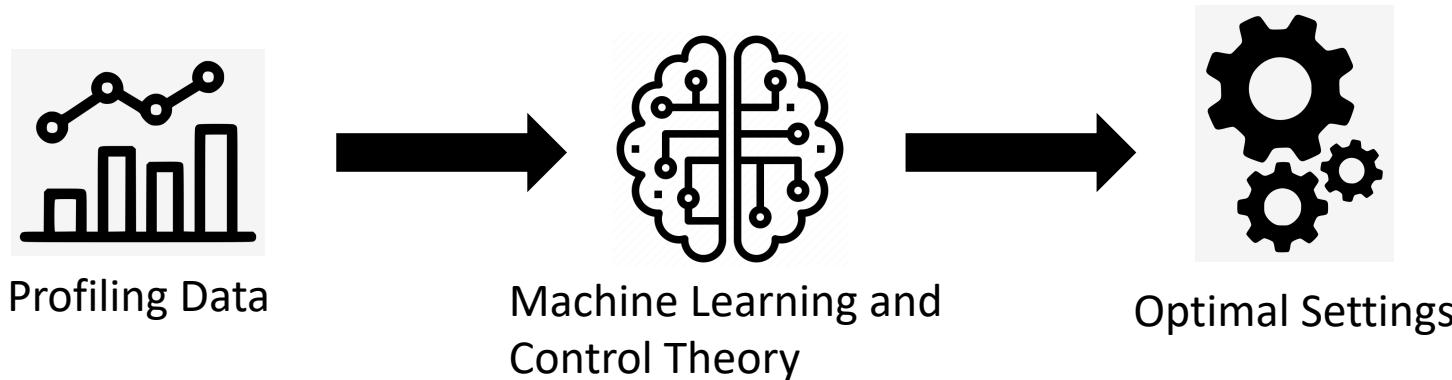
# Can we help?

Can we automatically answer ...

Does a configuration affect performance?

How does a configuration affect performance?

# Previous work ---- Auto-tuning



- Expensive training and profiling
- Not working if workload/environment changes at run time

此外，如果user想改一个模型之外的配置项，就没办法了

***How can we do better?***

# Our Key Insights

Dynamic behavior

Does a configuration  
affect performance?

How does a Performance-  
sensitive Configuration  
(PerfConf) affect performance?



Static program logic

Does a Performance  
Operation (PerfOp) depend  
on the configuration?

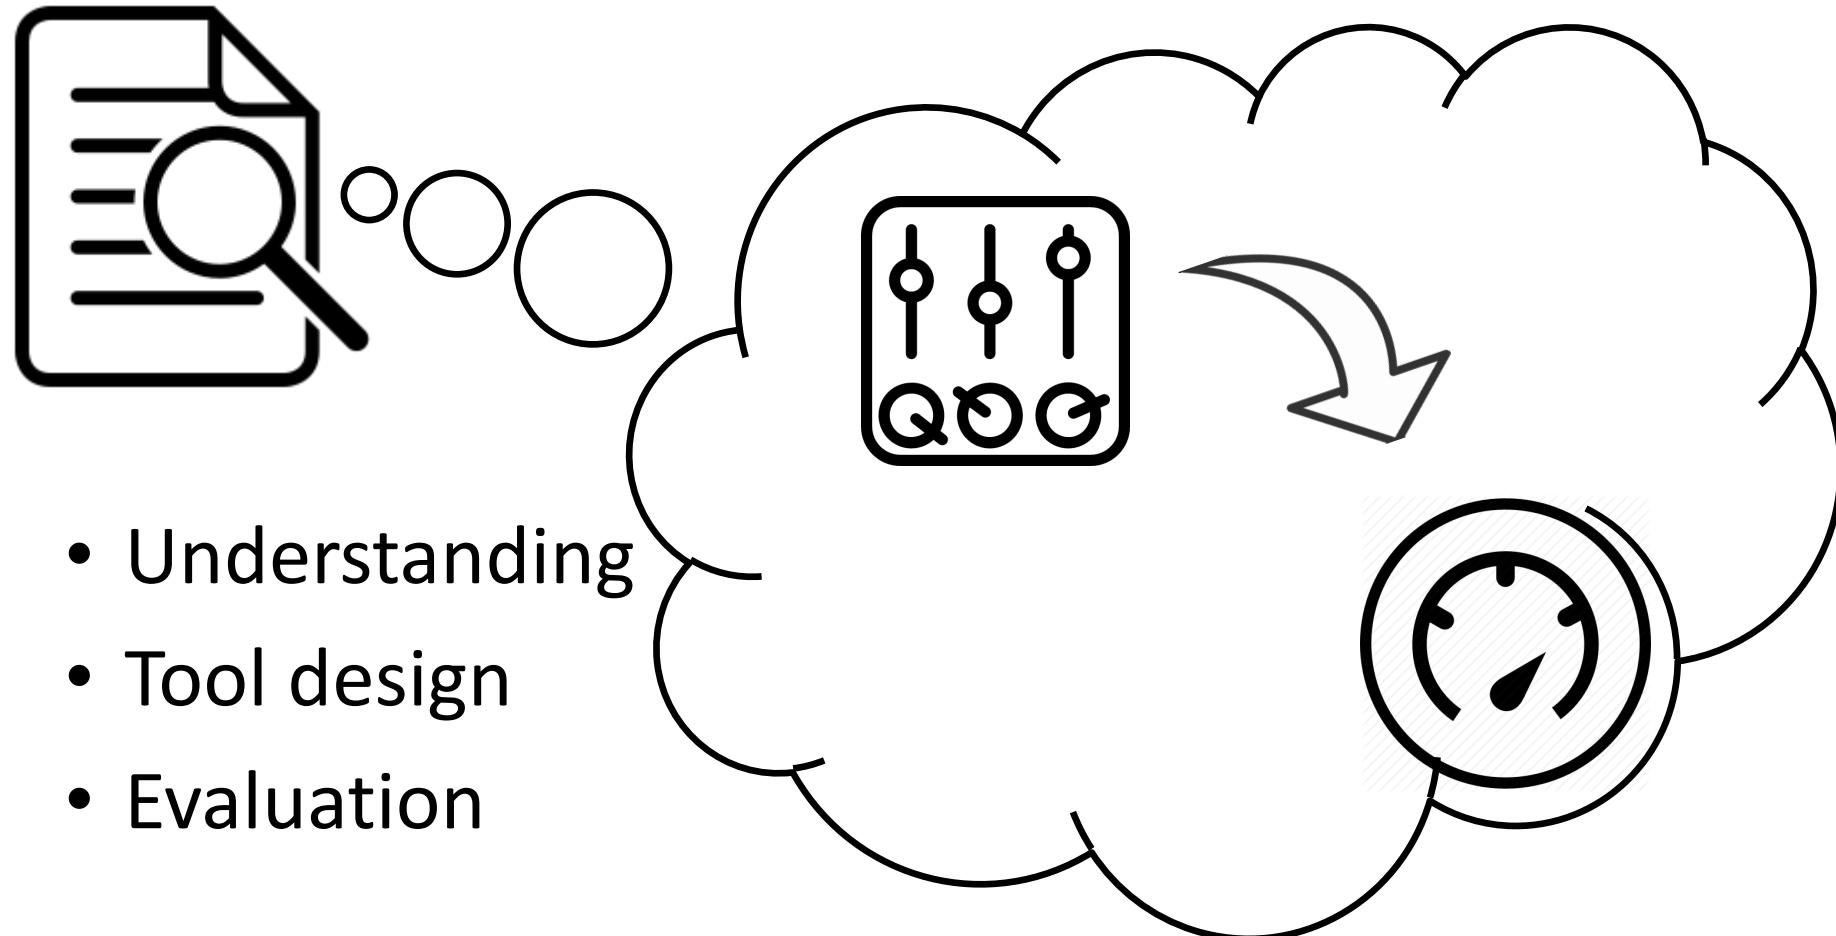
How does the  
PerfOp depend on  
the PerfConf?

**PerfConf**

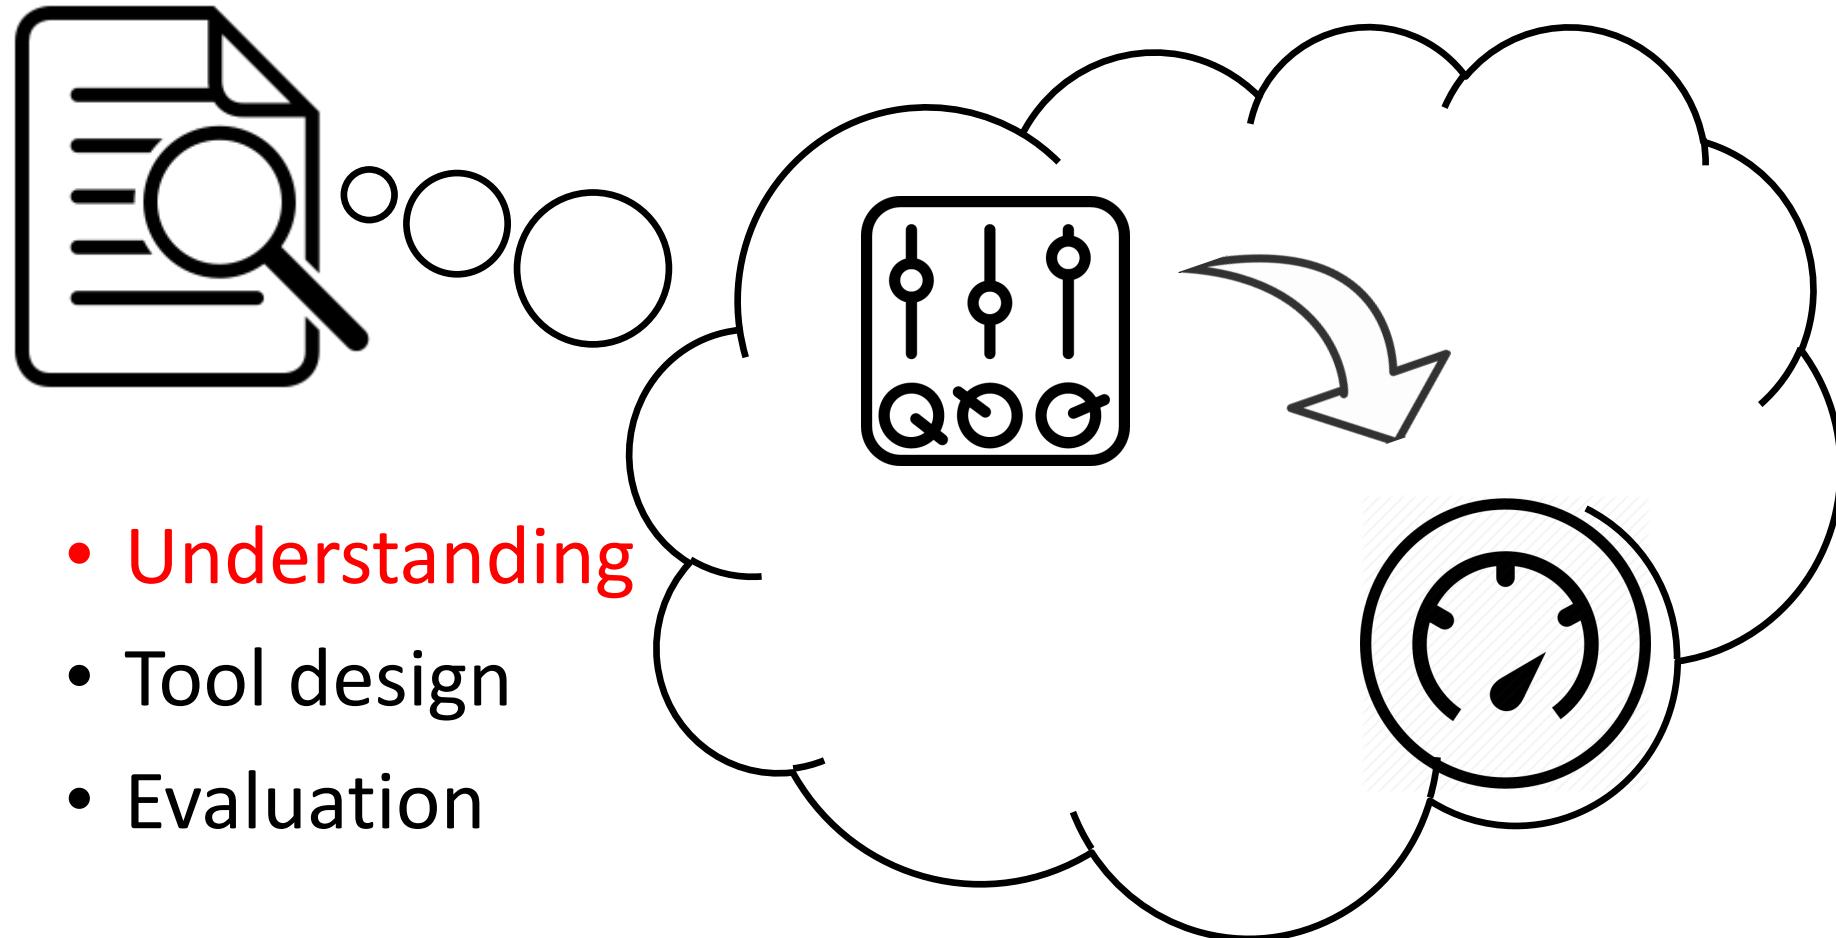
```
int sortmb = job.getInt("io.sort.mb");
int maxUsage = sortmb * 1024 * 1024;
buffer = new Byte[maxUsage];
```

**PerfOp**

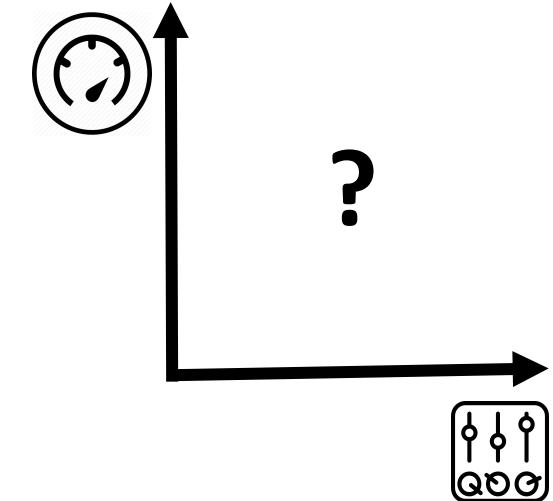
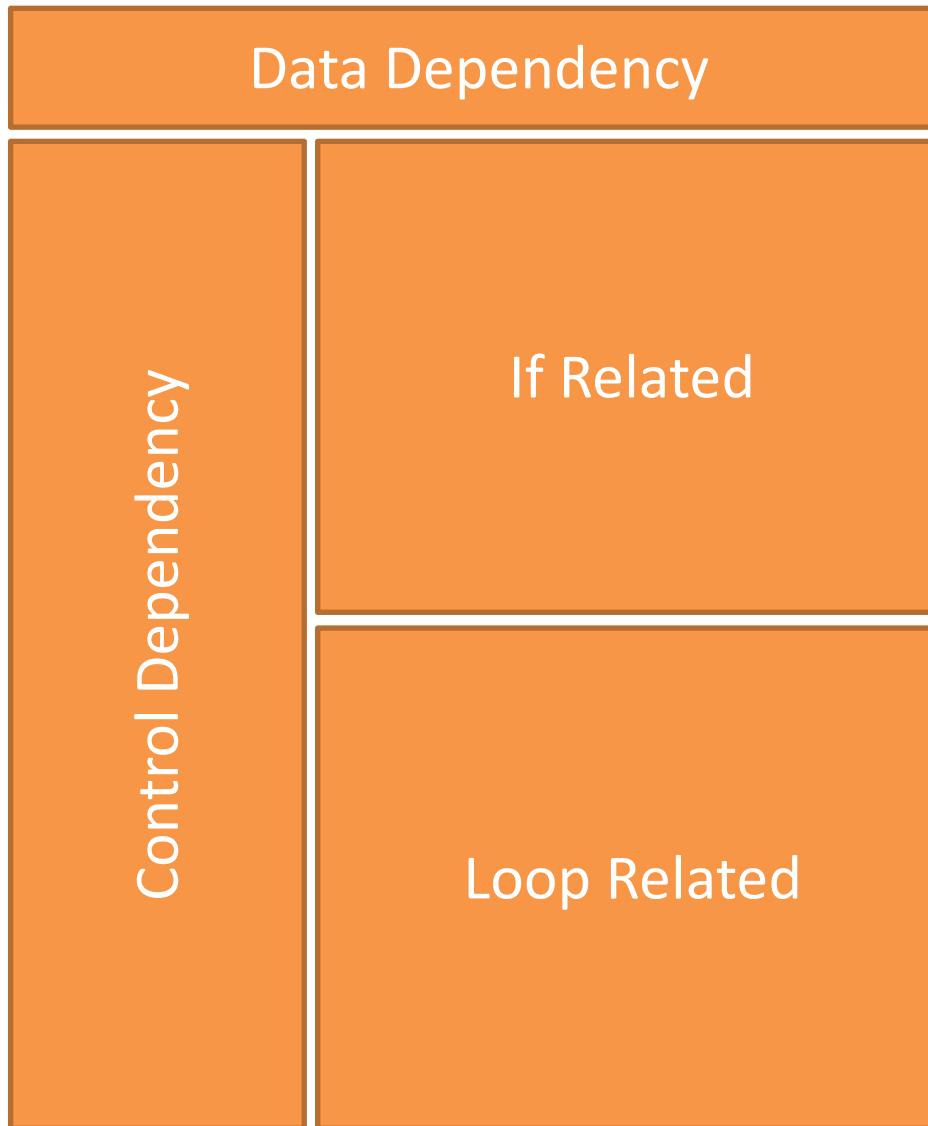
# *How to use program analysis to infer configurations' performance impact?*



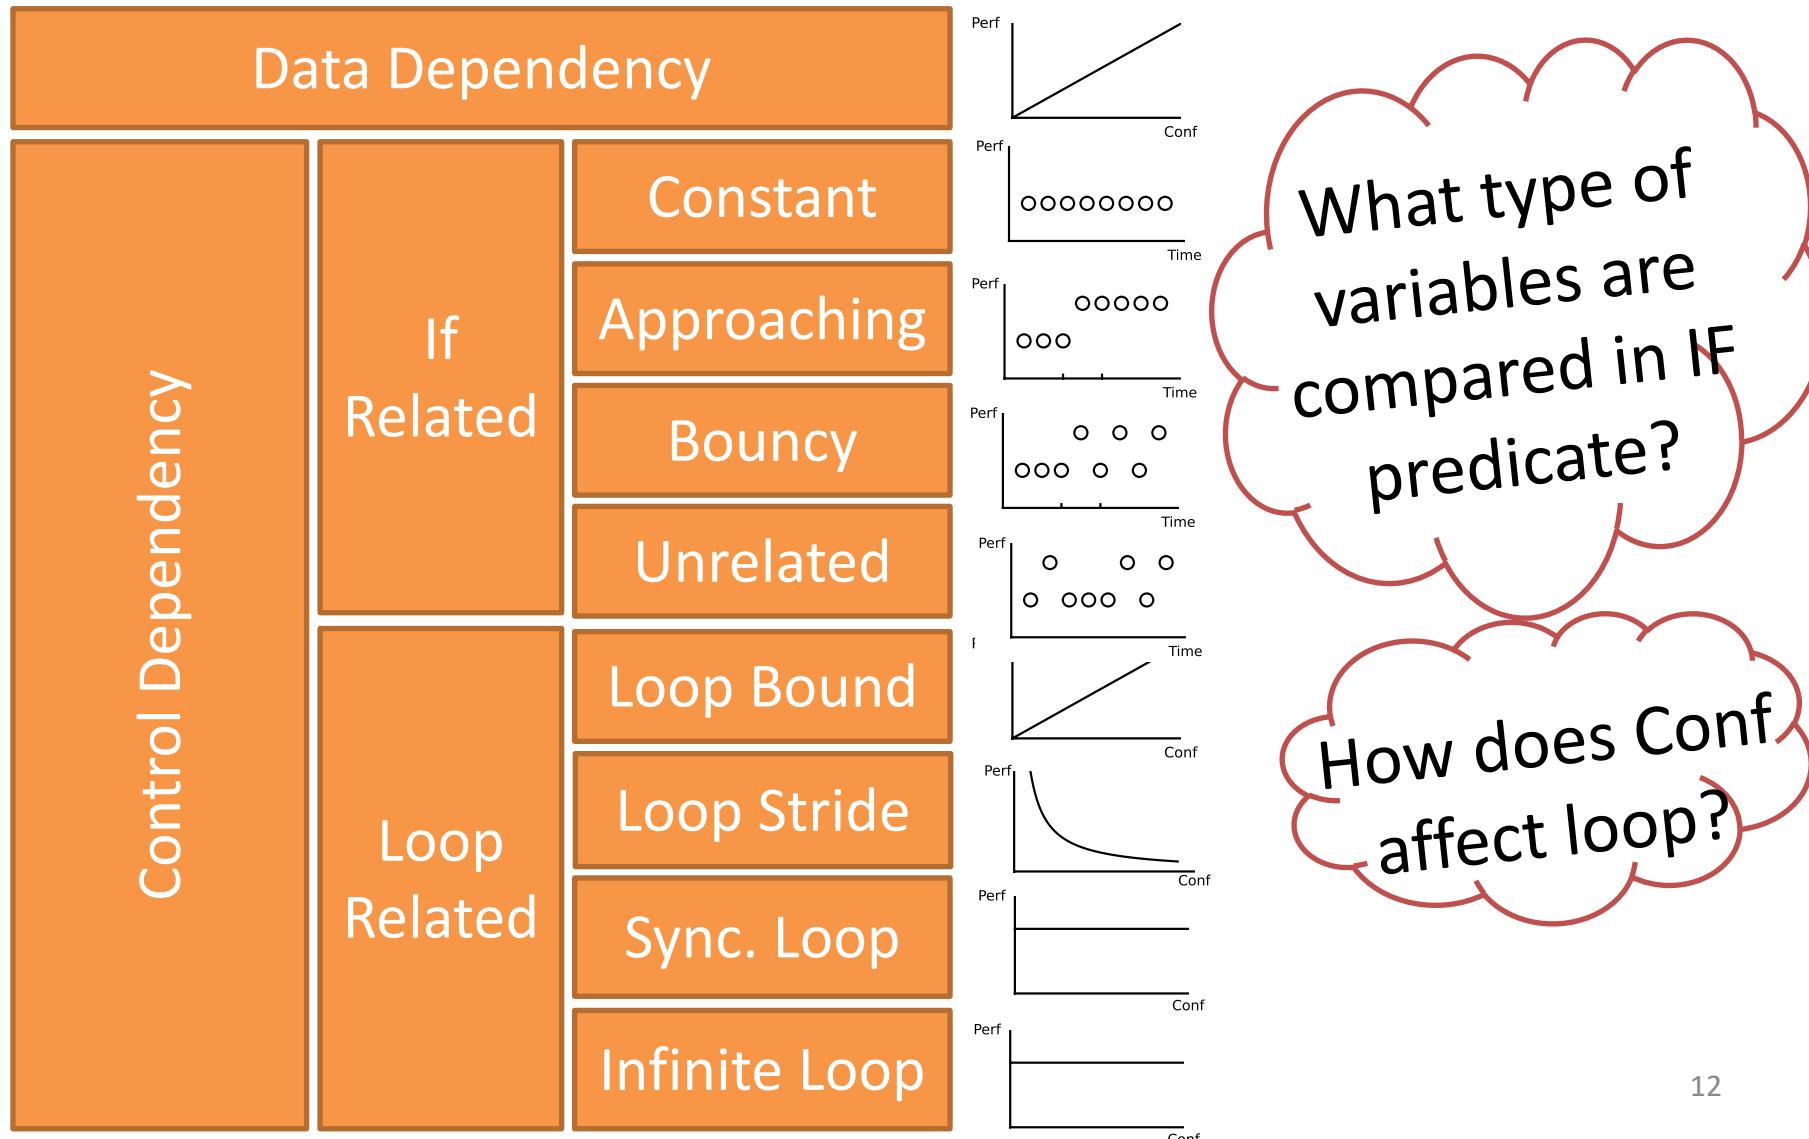
# *How to use program analysis to infer configurations' performance impact?*



# How can a Conf affect a Perf-Op?

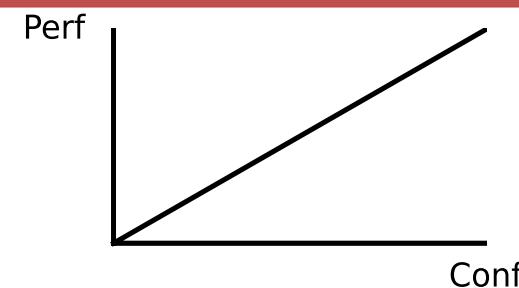


# How can a Conf affect performance?



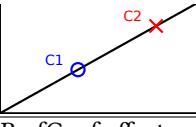
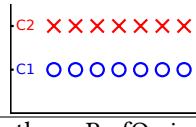
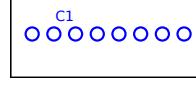
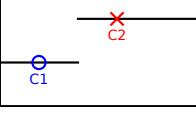
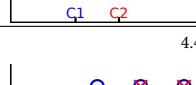
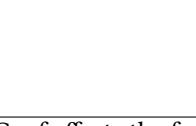
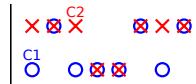
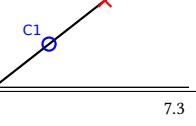
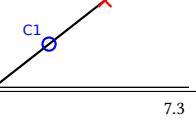
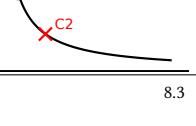
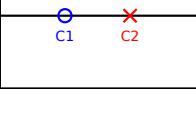
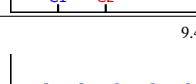
# Data Dependency

- Configuration affects the **impact of every instance** of PerfOp through parameters

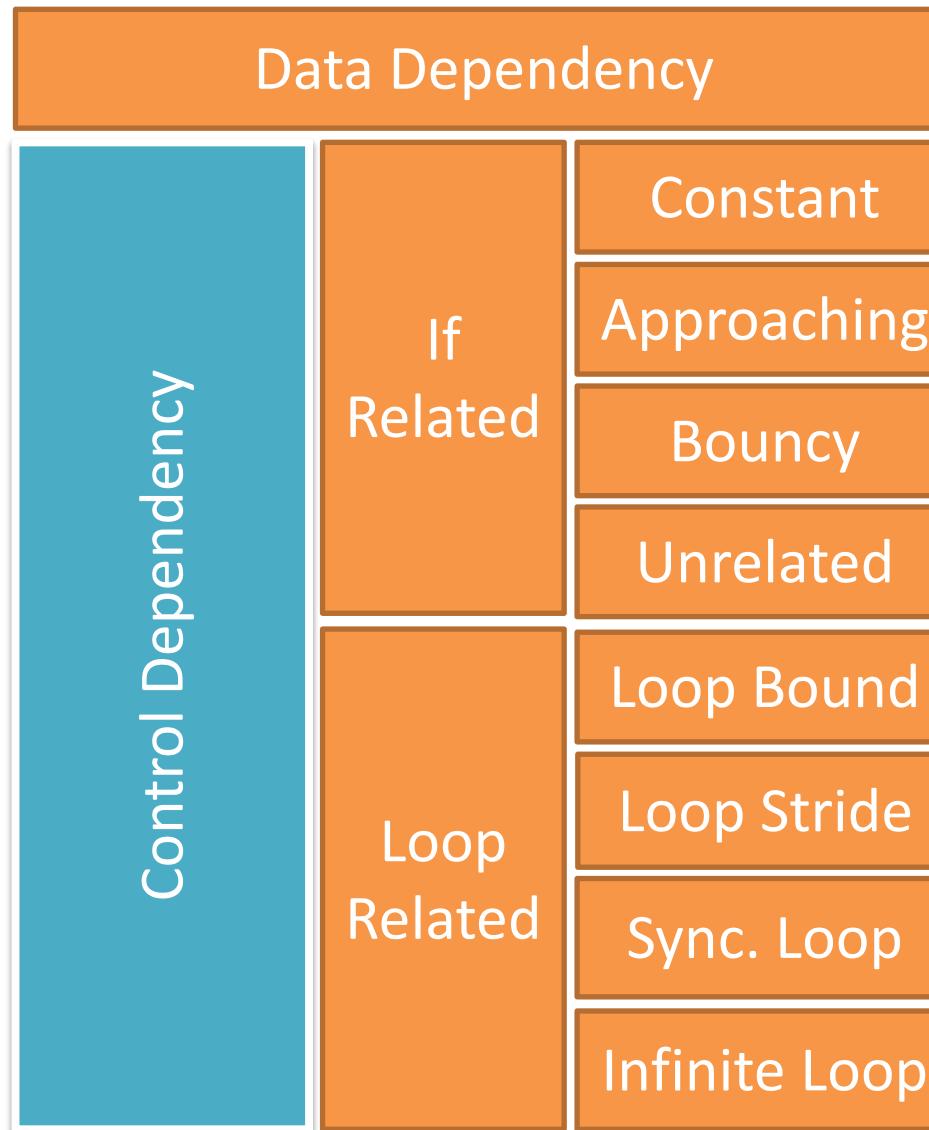
Code Example	Formula	Performance Graph
PerfOp(Conf)	$Performance = Conf$	

```
int sortmb = job.getInt("io.sort.mb");
int maxUsage = sortmb * 1024 * 1024;
buffer = new Byte[maxUsage];
```

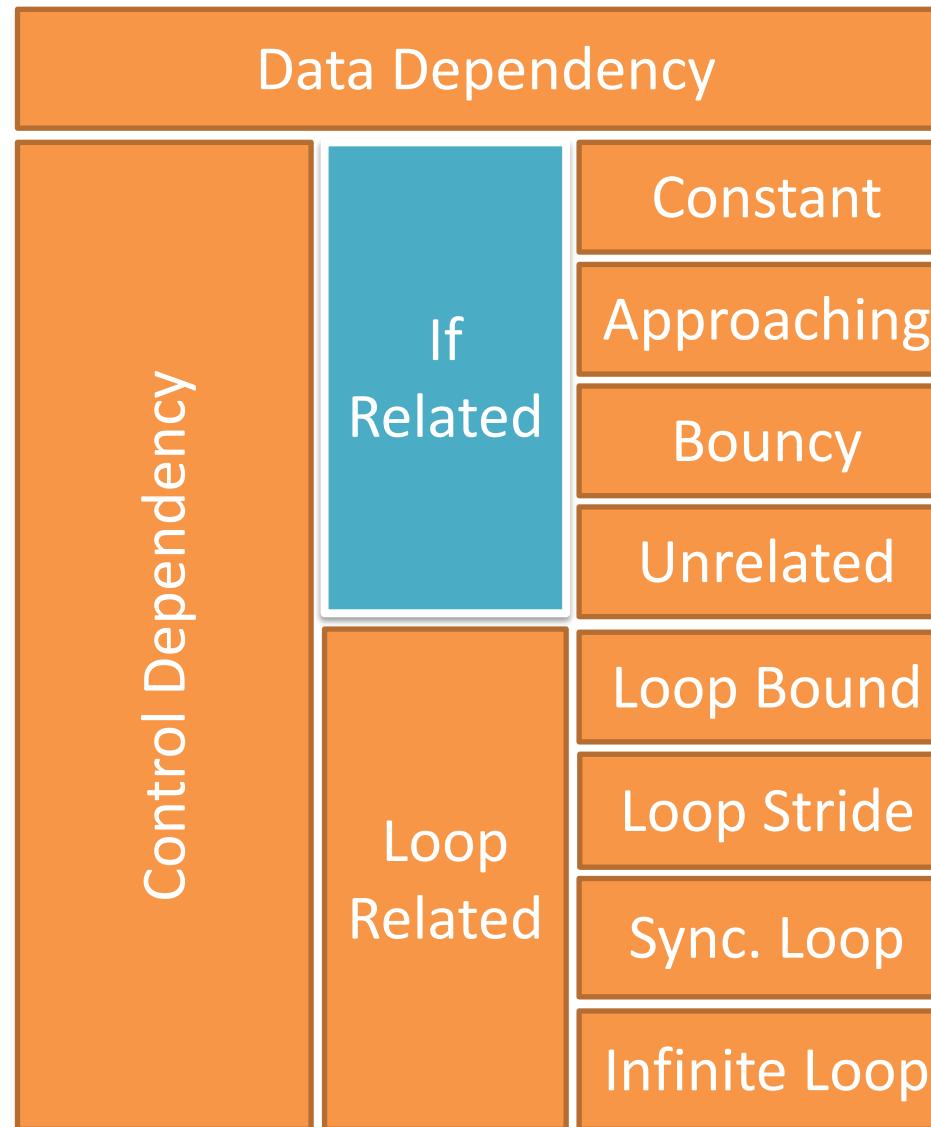
**Table 1. Configurations' Performance-Impact Taxonomy.** To ease the discussion, all the formulas, figures, and explanations in the table focus on the Toy Examples, where  $C$  is the value of a configuration and Perf refers to the memory-consumption contribution of one instance of the toy-example code snippet. The *Perf–Conf* figure in the 4th column depicts how the Perf (the y-axis) might change with the configuration setting (the x-axis); the *Perf–Time* figure in the 5th column depicts how the Perf (the y-axis) might change over time (the x-axis) under different configuration settings. In both columns' figures, we use C1 (blue dots) and C2 (red crosses) to represent two different configuration settings, and  $C_1 < C_2$ .

Pattern	Toy Example	Formula	Perf–Conf	Perf–Time	Explanation
<b>Data Dependency – PerfConf affects the amount of performance contribution from one PerfOp</b>					
1 Data	1.1 <code>new byte[C];</code>	1.2 $\text{Perf} = C$	1.3 	1.4 	$C$ affects the impact of one <code>new byte[]</code> linearly.
<b>Control Dependency (IF) – PerfConf affects whether a PerfOp is executed</b>					
2 w/ Con- stant	2.1 <code>if(V&lt;=C)     new byte[a]; else new byte[b];</code>	2.2 $V \leq C$	2.3 	2.4 	$C$ affects whether execute <code>new byte[a]</code> or <code>new byte[b]</code> statically.
3 w/ getting- closer- var	3.1 <code>if (V&lt;=C) {     V++; new byte[a]; } else     new byte[b];</code>	3.2 $V \leq C$	3.3 	3.4 	$C$ affects when to switch from <code>new byte[a]</code> to <code>new byte[b]</code> .
4 w/ back& forth var.	4.1 <code>if (V&lt;=C) {     V++; new byte[a]; } else {     V--; new byte[b]; }</code>	4.2 $V \leq C$	4.3 	4.4 	$C$ affect how quickly switch from <code>new byte[a]</code> to <code>new byte[b]</code> .
5 w/ un- related var.	5.1 <code>if(V&lt;=C)     new byte[a]; else new byte[b];</code>	5.2 $V \leq C$	5.3 	5.4 	$C$ affects the probability of <code>new byte[a]</code> executes over <code>new byte[b]</code> .
<b>Control Dependency (LOOP) – PerfConf affects the frequency/#-of-times a PerfOp executes</b>					
6 Regular loop bound	6.1 <code>for(; i&lt;C; i++) {     new byte[a]; }</code>	6.2 $\text{Perf} = a * C$	6.3 	6.4 	$C$ affects the number of <code>new byte[a]</code> via loop bound.
7 Regular loop stride	7.1 <code>for(; i&lt;N; i+=C) {     new byte[a]; }</code>	7.2 $\text{Perf} = a * N / C$	7.3 	7.4 	$C$ affects the number of <code>new byte[a]</code> via loop stride.
8 Sync loop	8.1 <code>while (i&lt;C) {     wait(); } new byte[a];</code>	8.2 $\text{Perf} = a$	8.3 	8.4 	$C$ affects the frequency of <code>new byte[a]</code> via synchronization loop.
9 Infinite loop	9.1 <code>new byte[a]; sleep(C); //in infinite loop</code>	9.2 $\text{Perf} = a$	9.3 	9.4 	$C$ affects the frequency of <code>new byte[a]</code> via a fixed interval.

# How can a Conf affect performance?

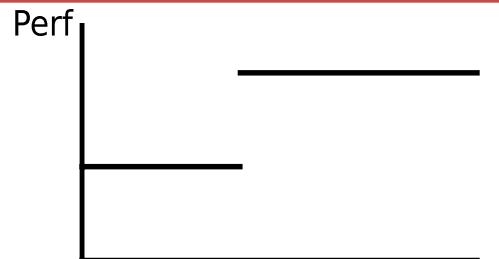


# How can a Conf affect performance?

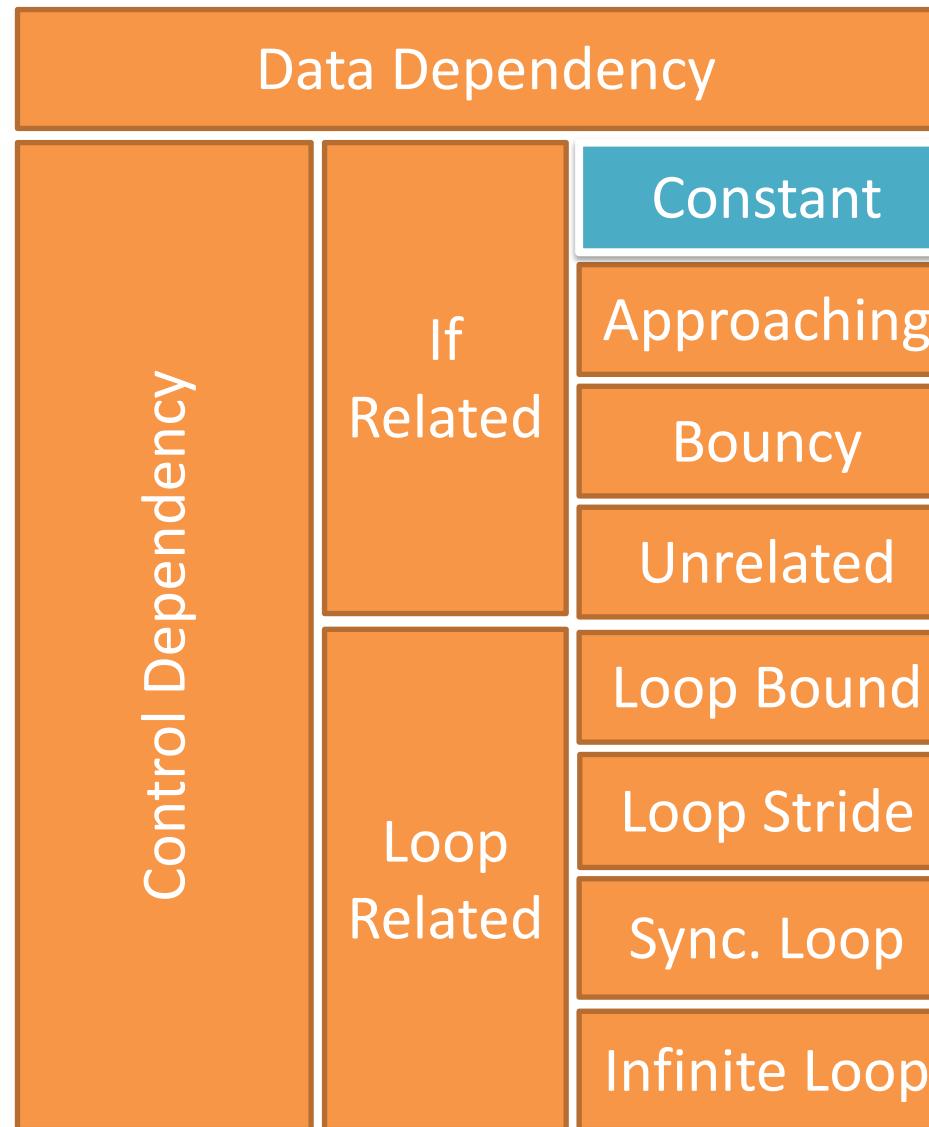


# If Related Patterns

- Conf affects whether the PerfOp is executed

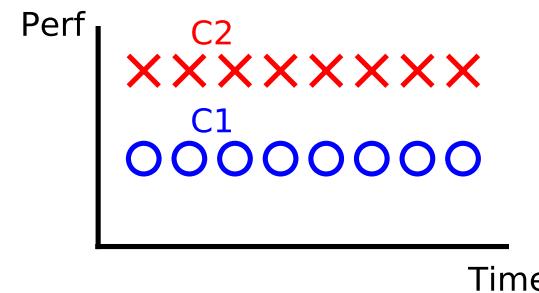
Code Example	Formula	Performance Graph
If $(V \leq C)$ { PerfOpA } else { PerfOpB }	Performance $= \begin{cases} a, & V \leq C \\ b, & V > C \end{cases}$	

# How can a Conf affect performance?

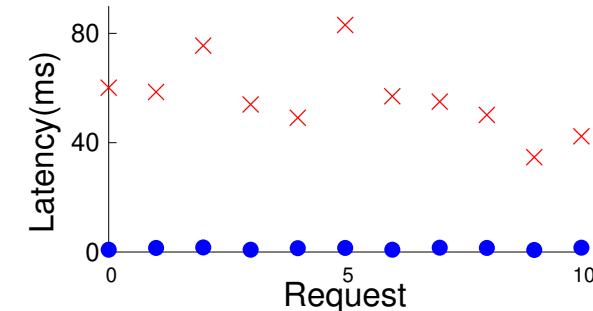


# Compared with Constant

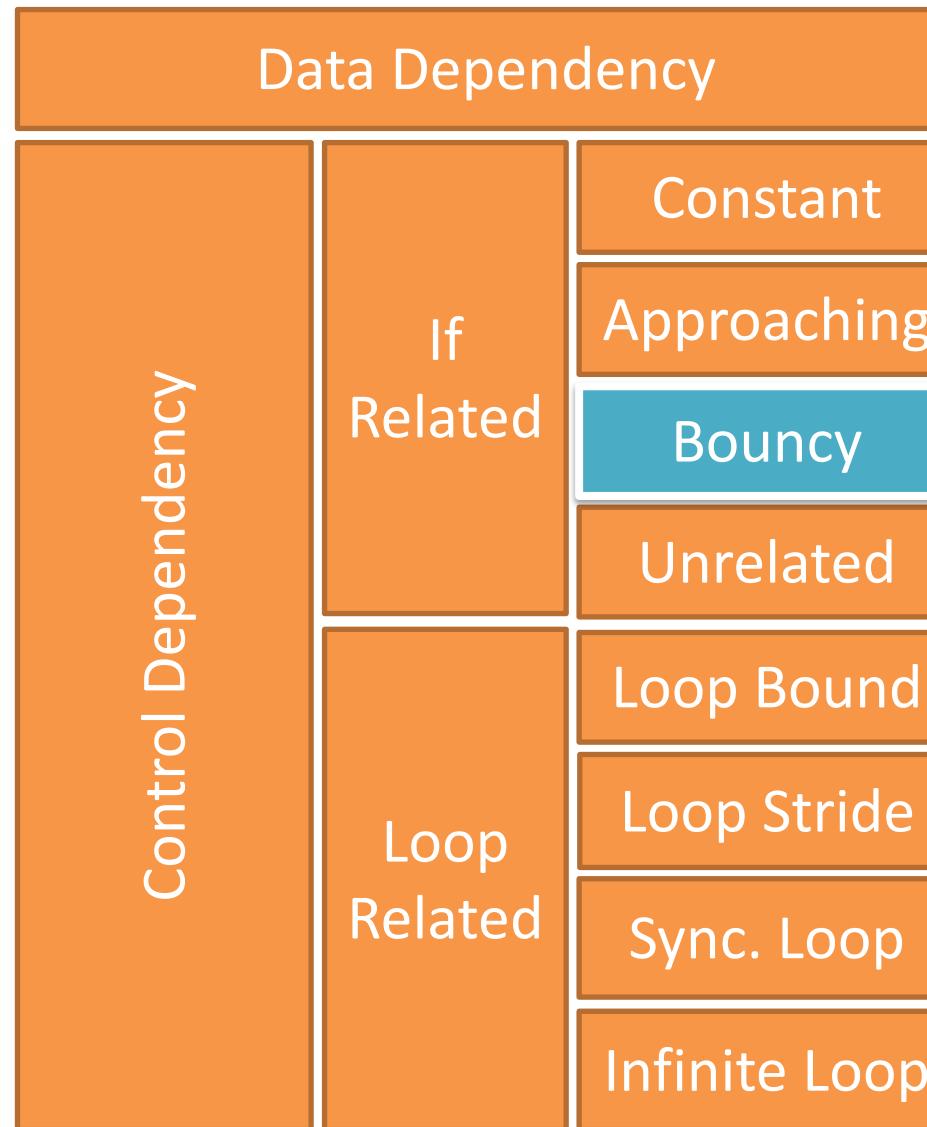
- The if-else decision does not change over time



```
if (maxFsObjects != 0) {  
    lock();  
}
```

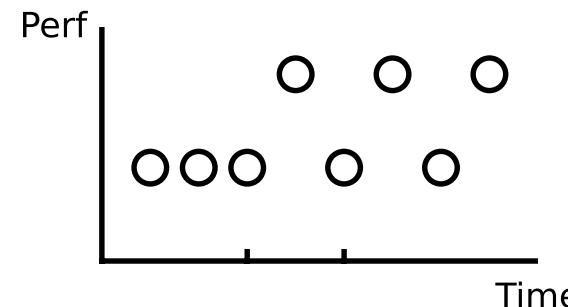


# How can a Conf affect performance?

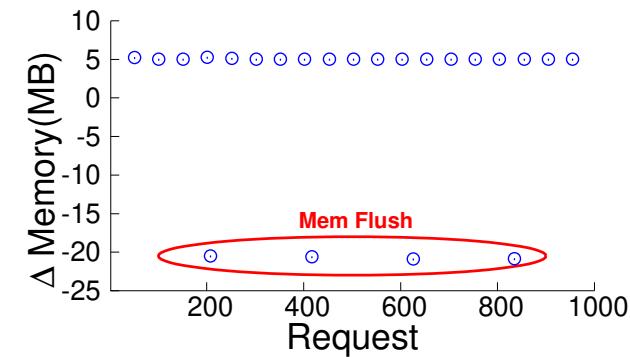


# Compared with Bouncy Variable

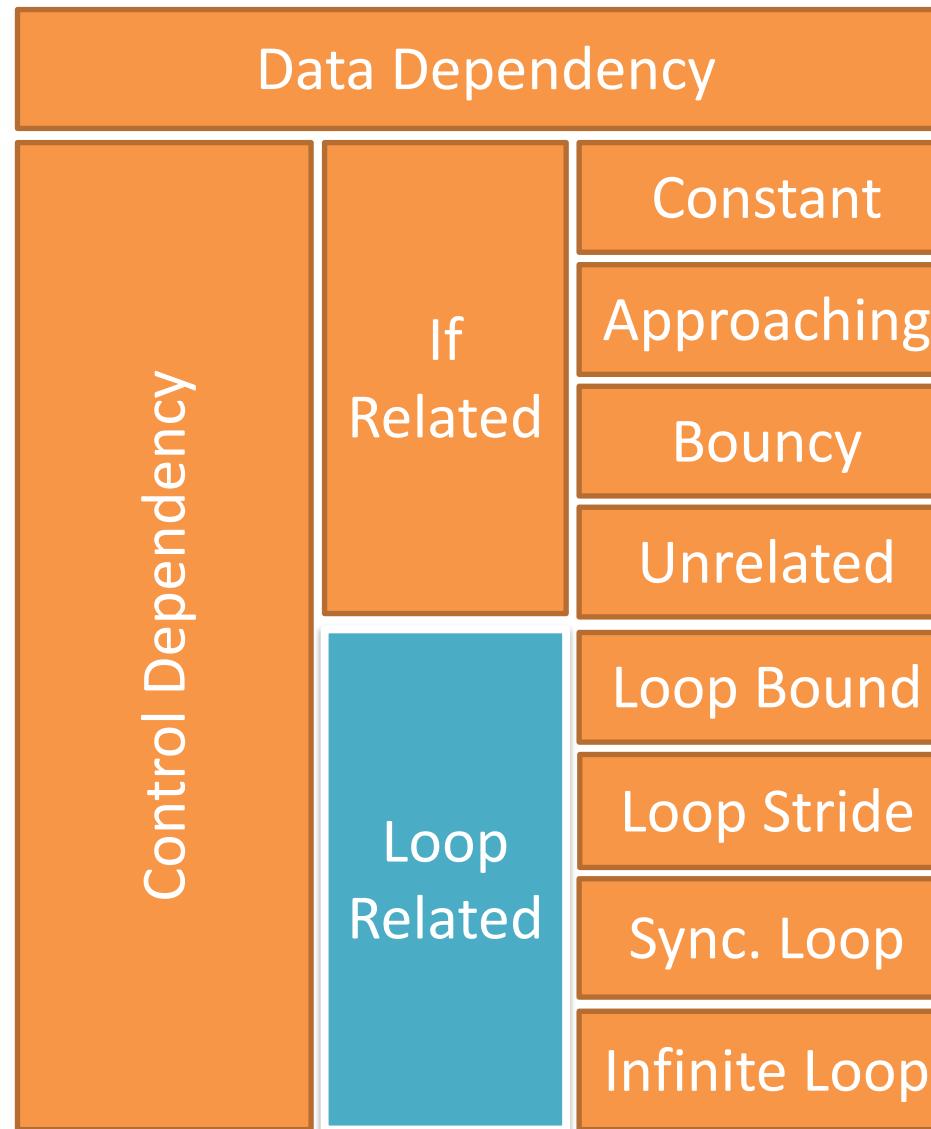
- The if-else decision keeps changing over time



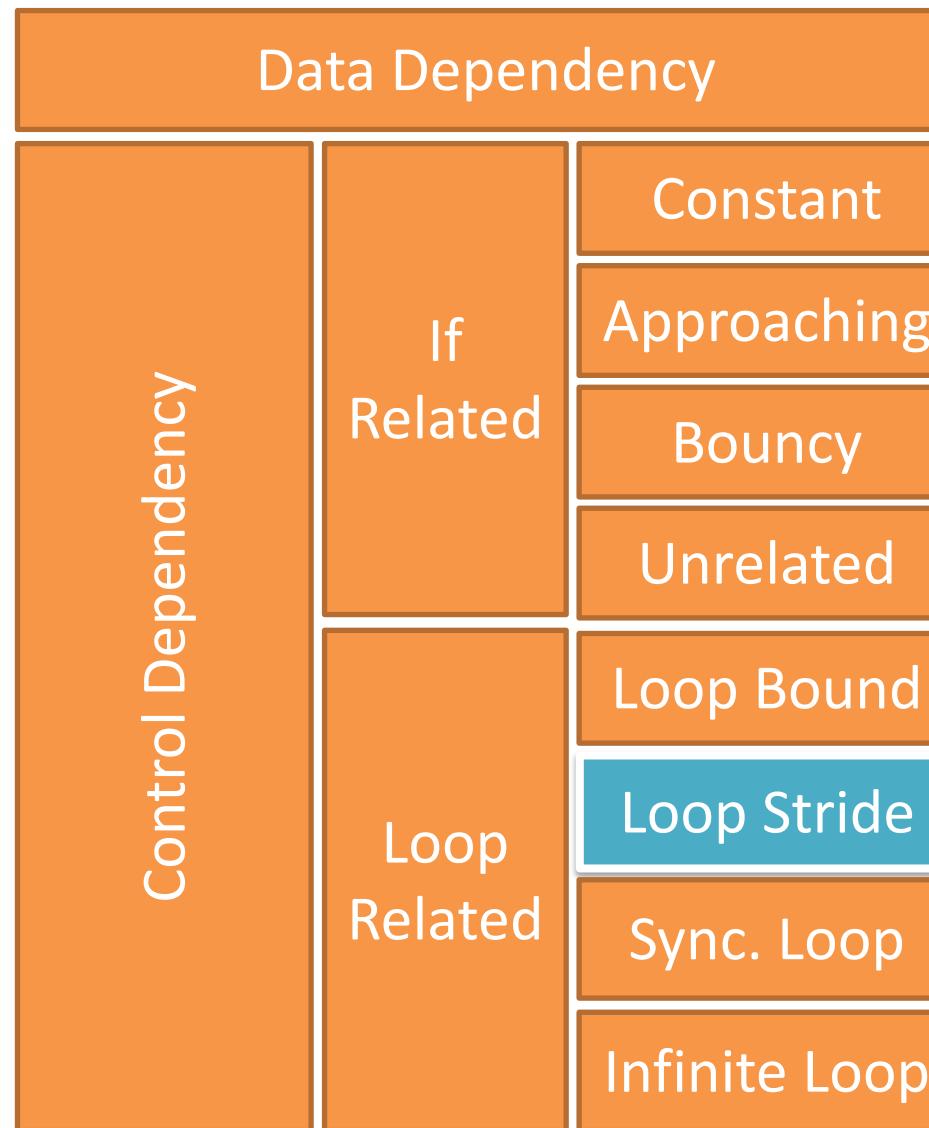
```
currentSize += put.heapSize();
writeBuffer.add(put);
if (currentSize > bufferSize) {
    writeBuffer clear();
    currentSize = 0;
}
```



# How can a Conf affect performance?

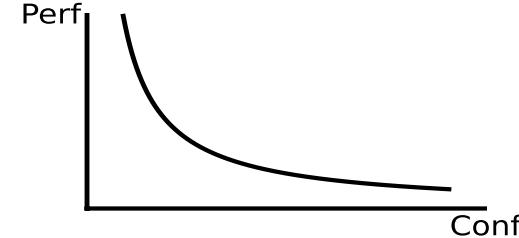


# How can a Conf affect performance?

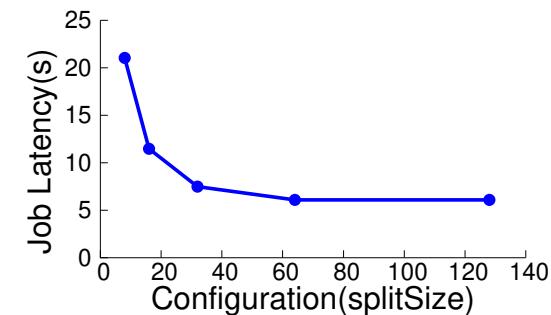


# Affect Loop Stride

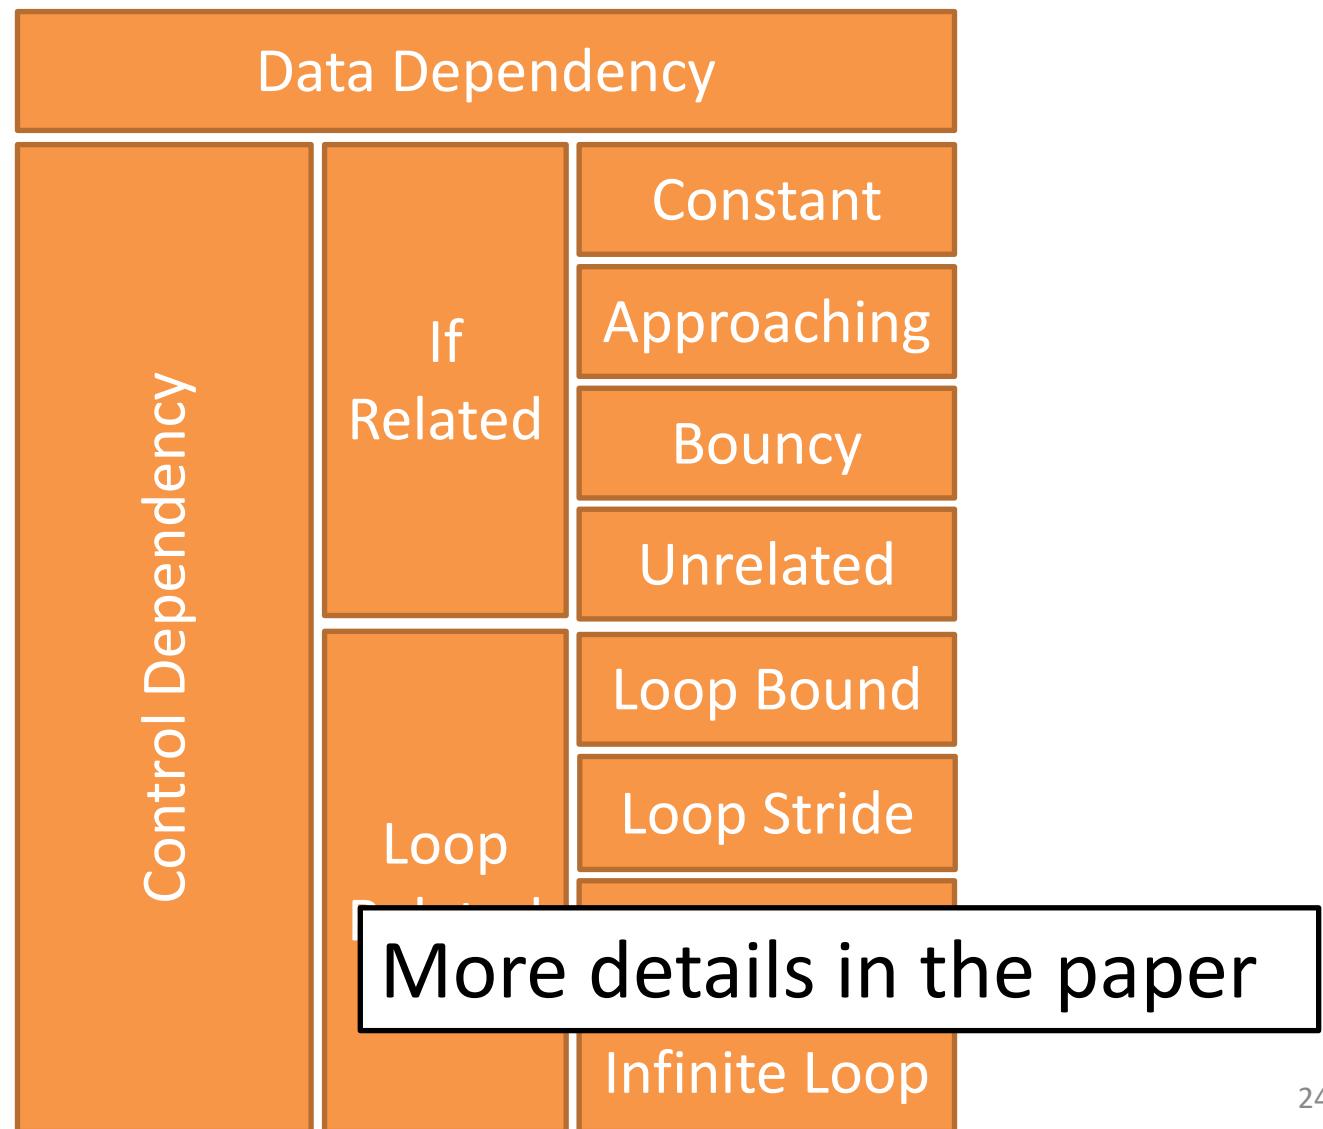
- Conf used as a loop stride in the loop-exit condition

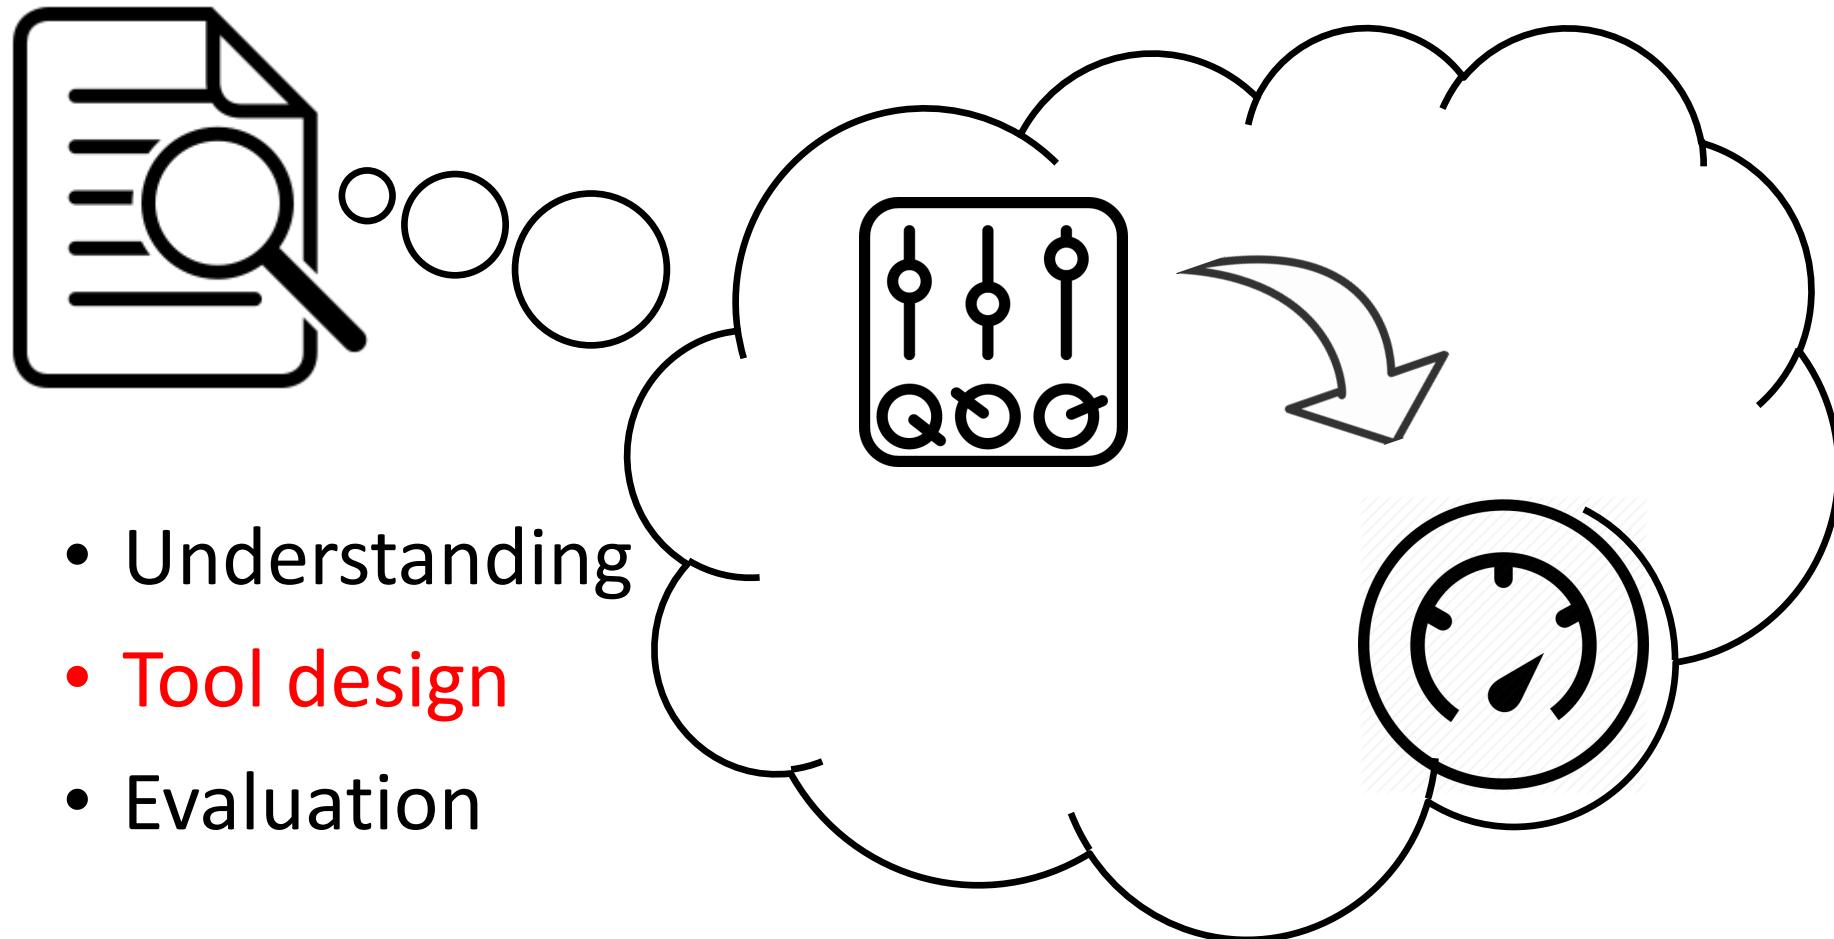
Code Example	Formula	Performance Graph
<pre>for ( ; i &lt; N; i+=Conf) {     PerfOp(); }</pre>	$\text{Performance} = aN/\text{Conf}$	

```
while (bytesRemaining > 0) {
    splits.add(makeSplit());
    bytesRemaining -= splitSize;
}
```

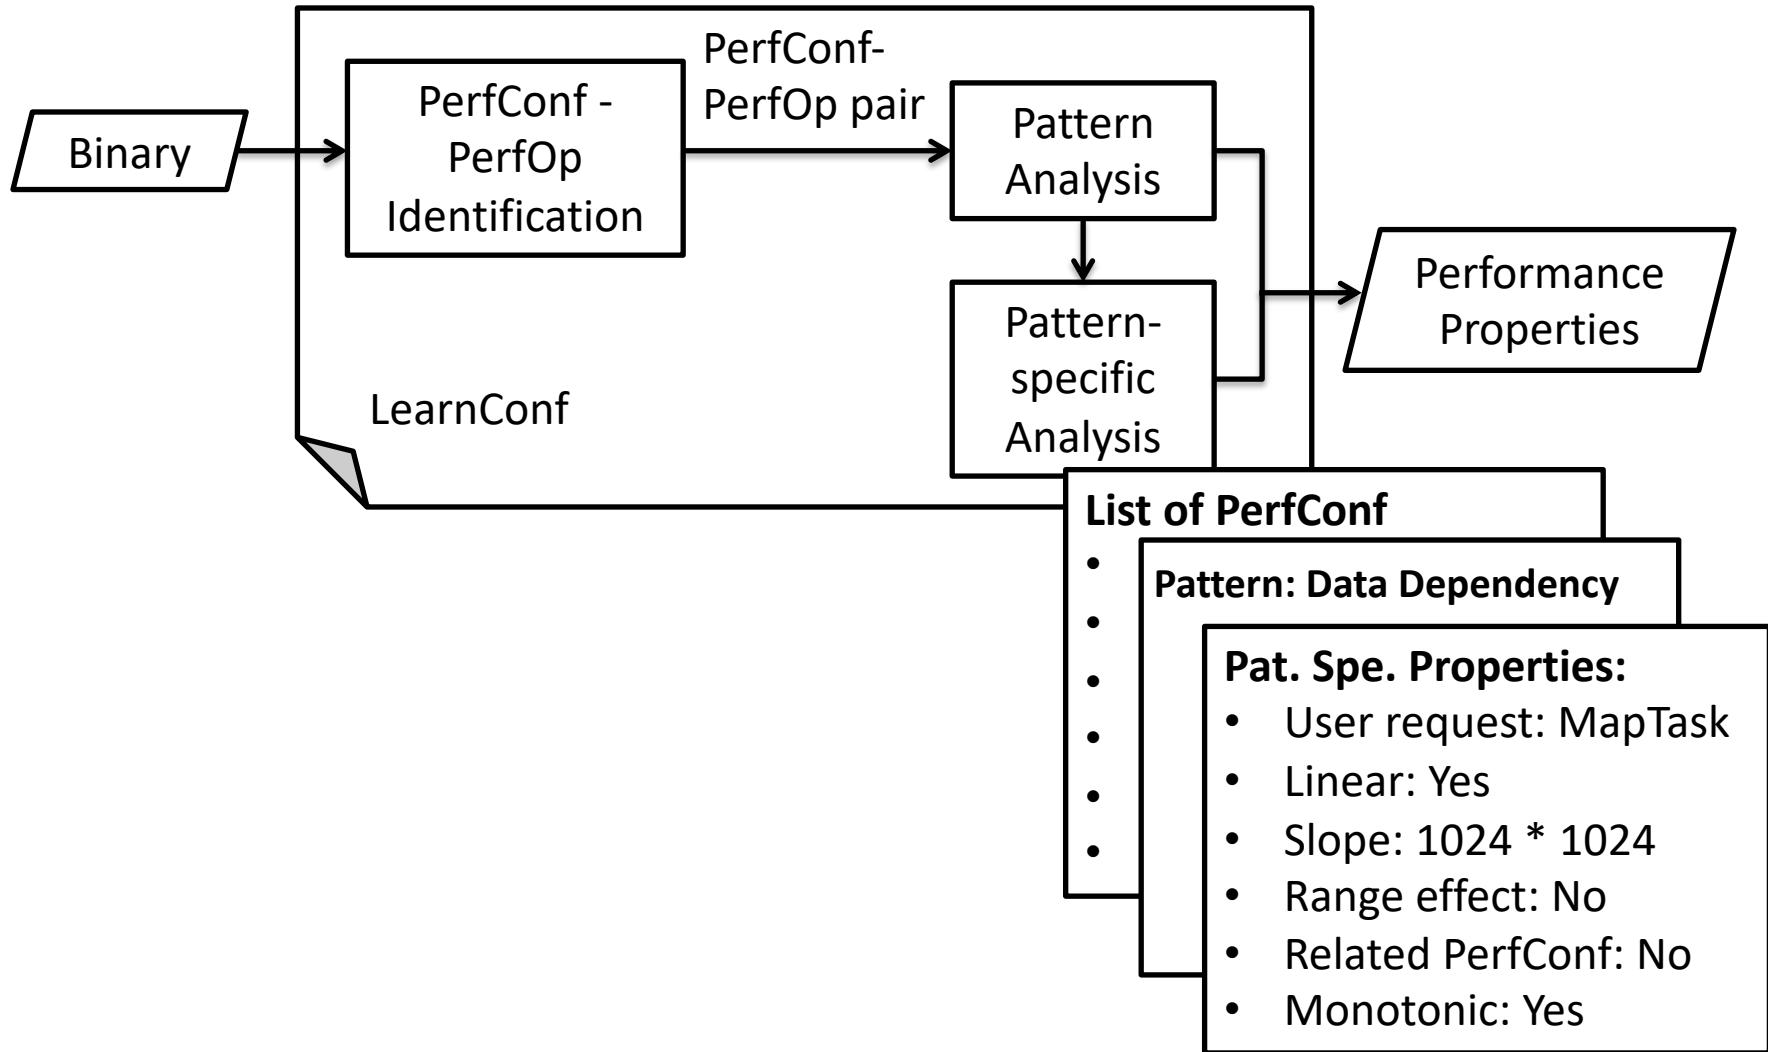


# How can a Conf affect performance?





# LearnConf Overview



# Identify Configuration Variable

- Identify configuration-loading API
  - Add return var. to configuration variable set
- Track data-dependence chain
  - Tag more variables as configuration variables

Configuration variable

configuration-loading API

```
int sortmb = job.getInt("io.sort.mb");
int maxUsage = sortmb * 1024 * 1024;
buffer = new Byte[maxUsage];
```

# Identify PerfOps

- Latency related
  - Sleep(), lock(), IO, etc.
- Memory related
  - new byte[], List.add(), etc.

```
int sortmb = job.getInt("io.sort.mb");
int maxUsage = sortmb * 1024 * 1024;
buffer = new Byte[maxUsage];
```

Memory Intensive Operation

# Identify PerfConf

If a PerfOp depends on the Configuration Variable, ...

```
int sortmb = job.getInt("io.sort.mb");
int maxUsage = sortmb * 1024 * 1024;
buffer = new Byte[maxUsage];
```

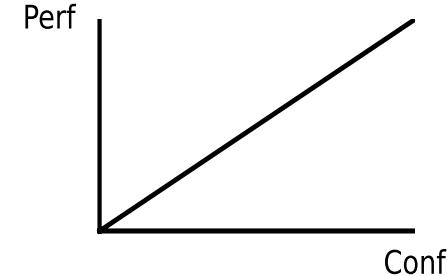
- List of PerfConf**
- io.sort.mb
  - ...

# Categorize PerfConf-PerfOp dependency

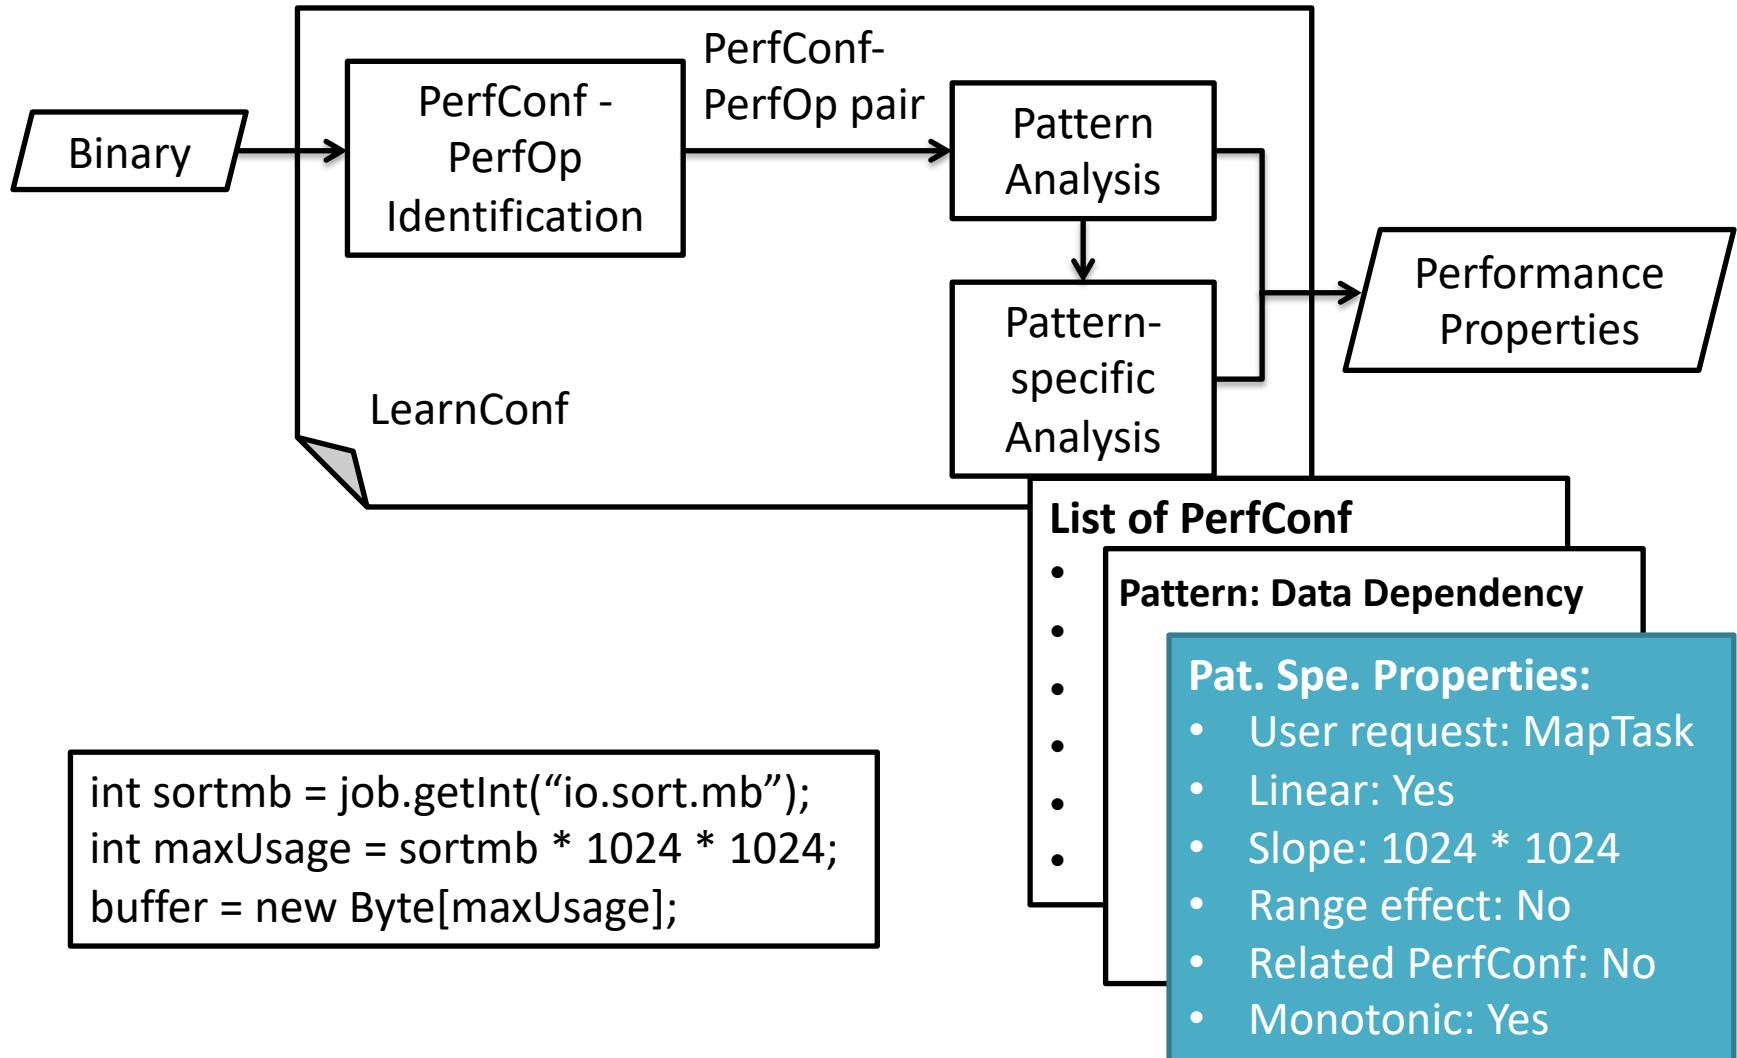
- Data Dependency Pattern
  - Conf used in the parameter of the PerfOp
- If Pattern
  - Conf used in an if-predicate
- Loop Pattern
  - Conf used in a loop-exit condition

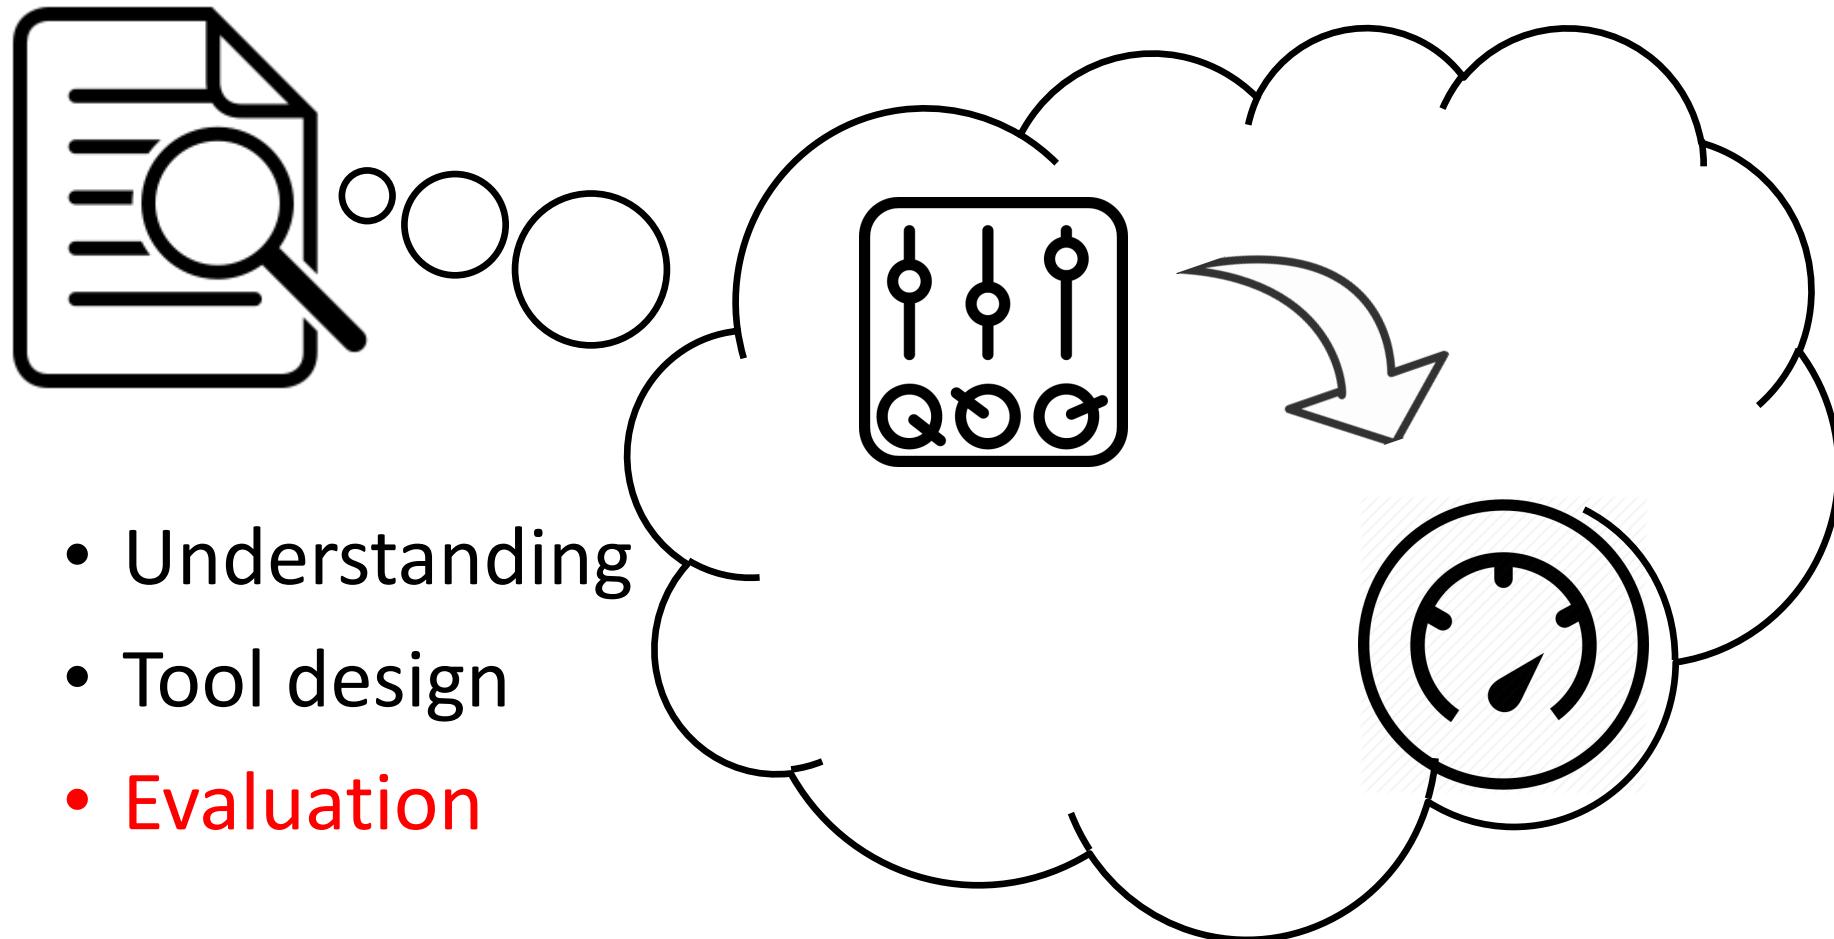
```
int sortmb = job.getInt("io.sort.mb");
int maxUsage = sortmb * 1024 * 1024;
buffer = new Byte[maxUsage];
```

**Pattern: Data Dependency**



# Pattern-Specific Analysis





# Methodology

- Benchmarks
  - Four widely used distributed systems
  - Each contains around 100~150 configurations



# Identify Correct PerfConf

UNION of tutorials and papers

- Correctly identify 60 out of 71 true PerfConfs
- 9 false positives
- 4 true PerfConfs **not** in previous work that can lead to OOM or timeout failures!  
workload, memory consumption

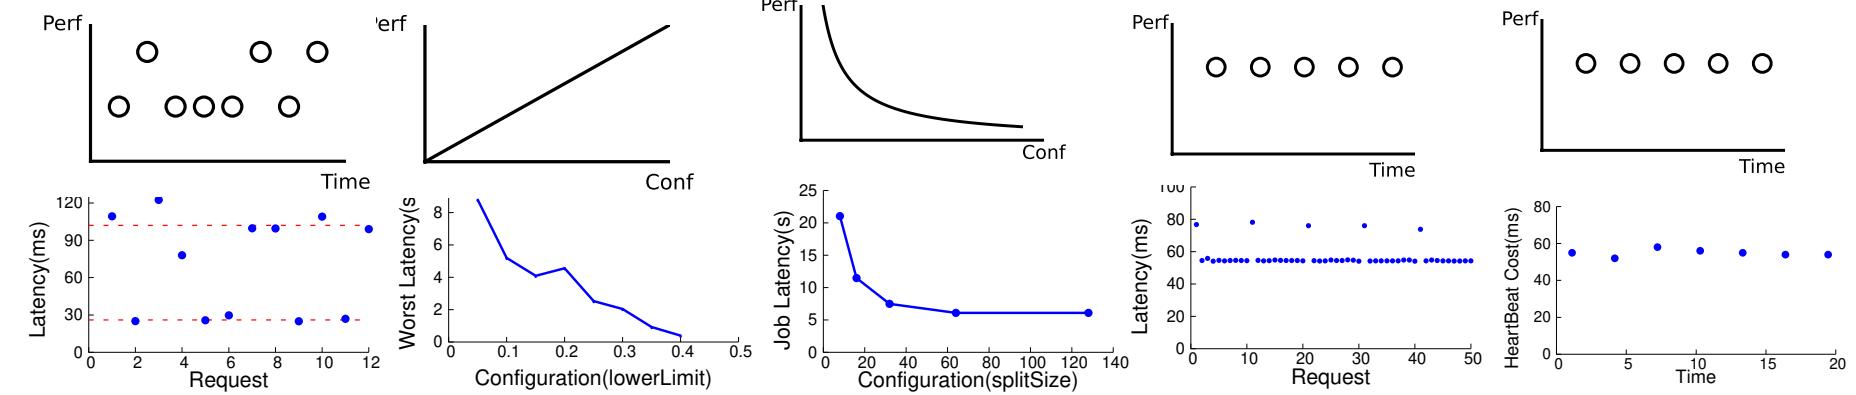
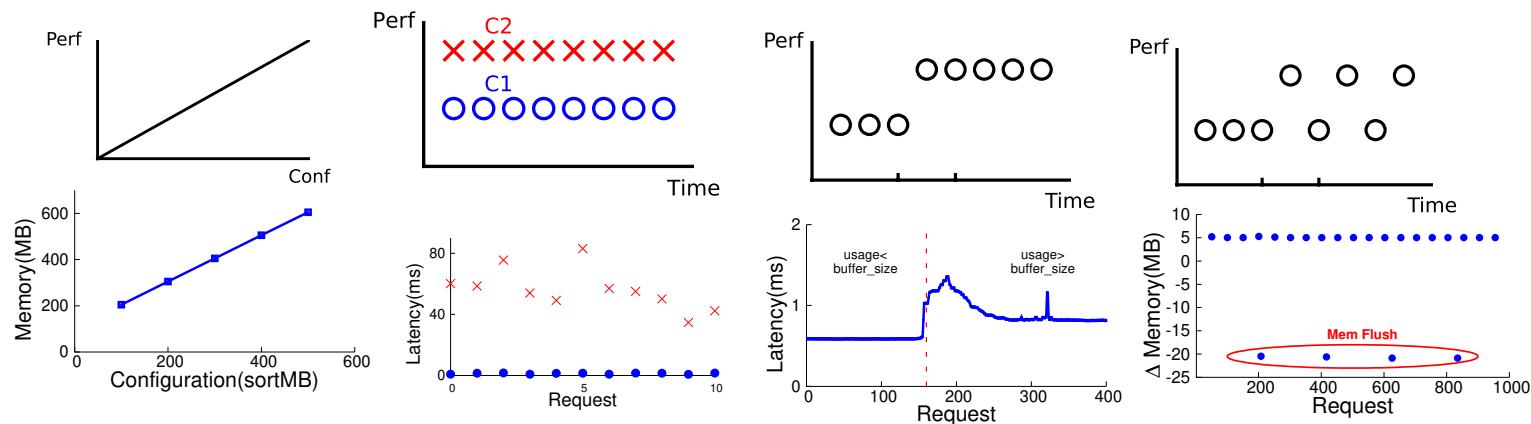
performance impact

eliminated by software mechanism

complicated control dependencies, IPC...

	Identified	False Positive	False Negative
MapReduce	16	1	7
HBase	19	1	2
HDFS	13	5	1
Cassandra	21	2	1
Total	69	9	11

# Identify Correct Pattern

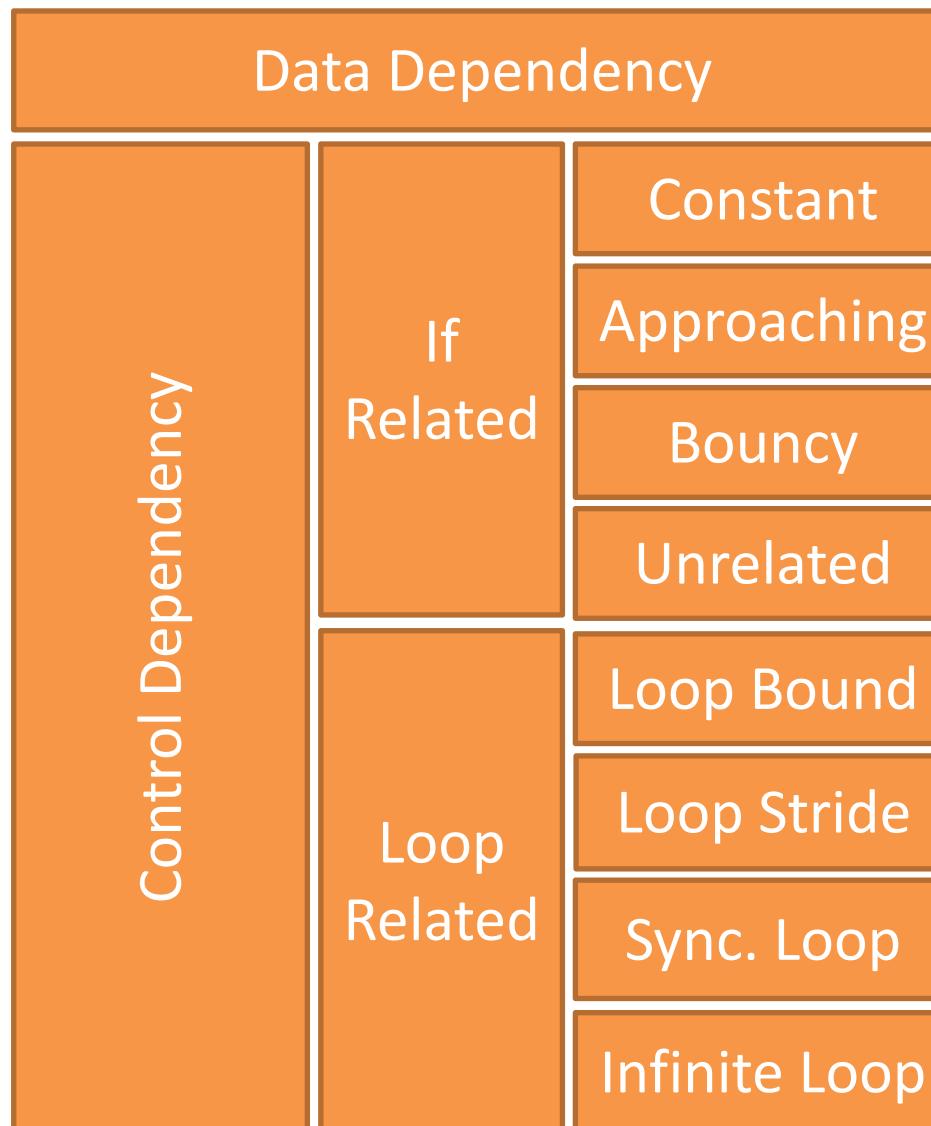


# More Result

- Input Analysis
- Slope Analysis
- Configuration Setting Range Analysis
- Configuration Relation Analysis
- Monotonicity Analysis
- Applying LearnConf for Performance Tuning

More results in the paper

# Conclusion



**Thanks**

Chi Li  
liche@uchicago.edu